University of Glasgow | School of Computing Science

Honours Individual Project Dissertation

TriBench: A Framework for Benchmarking SQL Workloads on Distributed Data Lakehouses

2736414Y

March 26, 2026

Abstract

The Apache Trino and Iceberg ecosystem lacks a reproducible, systematic benchmarking framework: existing tools require significant manual configuration, lack integrated hardware monitoring, or cannot port across deployment environments without modification. TriBench is a cross-platform benchmarking framework that addresses this gap by providing declarative experiment definitions, concurrent hardware and per-pod resource monitoring, and persistent result storage in a portable, self-contained bundle. The same experiment specification runs unchanged on a local Docker deployment or a multi-node Kubernetes cluster by switching a single configuration profile. Evaluated using TPC-H and TPC-DS across three environments (local Docker, a self-managed Kubernetes cluster, and Google Kubernetes Engine), TriBench produced reproducible query timings with full telemetry at scale factors up to SF 100, providing the end-to-end deployment, orchestration, and benchmarking capability that the Trino-Iceberg lakehouse ecosystem has lacked.

Acknowledgements

Education Use Consent

I hereby grant my permission for this project to be stored, distributed and shown to other University of Glasgow students and staff for educational purposes. **Please note that you are under no obligation to sign this declaration, but doing so would help future students.**

Signature: Lin Yuan Date: 20 March 2026

Contents

1 Introduction **1**
- 1.1 Motivation 1
- 1.2 Aims 1
- 1.3 Outline 2

2 Background **3**
- 2.1 Data Lakehouse Architecture 3
- 2.2 Benchmarking for Distributed SQL Systems 4
- 2.3 TPC Benchmark Standards 4
- 2.4 Reproducibility in Systems Research 5

3 Related Work **6**
- 3.1 Benchmarking Frameworks for Distributed Systems 6
  - 3.1.1 PEEL 6
  - 3.1.2 LST-Bench 7
  - 3.1.3 DiPerf 7
- 3.2 Lakehouse Technologies and Container Orchestration 7
  - 3.2.1 Apache Iceberg 8
  - 3.2.2 Apache Trino 8
  - 3.2.3 Kubernetes for Stateful Data Systems 9

4 Analysis and Requirements **10**
- 4.1 User Scenarios 10
- 4.2 Requirements Set 10
  - 4.2.1 Functional Requirements 10
  - 4.2.2 Non-Functional Requirements 11
- 4.3 Summary 11

5 Design **12**
- 5.1 Architectural Overview 12
- 5.2 Open Lakehouse Stack 13
- 5.3 CLI Layer 13
- 5.4 Bundle Abstraction 14
- 5.5 Configuration Layer 14
- 5.6 Orchestration Layer 15
  - 5.6.1 Experiment Definitions 15
  - 5.6.2 Execution Lifecycle 16

5.7 Monitoring System 17
5.8 Systems Layer 17
5.9 Dataset Loading 18
5.10 Result Storage 18
5.11 Reproducibility Mapping 19
5.12 Summary 19

6 Implementation 20
6.1 Development Approach 20
6.1.1 Language and Tooling 20
6.1.2 Version Control and Testing 20
6.2 CLI Implementation 20
6.3 Bundle Implementation 21
6.4 Configuration Loading 21
6.5 Dataset Loading 22
6.5.1 Built-in Loading 22
6.5.2 Universal Loader 22
6.6 Experiment Execution 23
6.6.1 Single Experiment 23
6.6.2 Suite Execution 24
6.6.3 Parallel Query Execution 24
6.7 Concurrent Monitoring 24
6.7.1 Threading Model 24
6.7.2 Host Resource Collector 25
6.7.3 Kubernetes Pod Collector 25
6.8 Deployment Backends 25
6.8.1 Docker Backend 25
6.8.2 Kubernetes Backend 25
6.9 Result Storage 26
6.9.1 Persistence 26
6.9.2 Export 26
6.9.3 Analysis Commands 26
6.10 Summary 27

7 Evaluation 28
7.1 Unit Testing 28
7.2 Evaluation Methodology 28
7.3 Experimental Setup 29
7.3.1 Local Environment (Level 1) 29
7.3.2 University of Glasgow Kubernetes Environment (Level 2) 29
7.3.3 GCP GKE Environment (Level 3) 29
7.4 Level 1: Local Docker Baseline 30
7.4.1 Deployment and Configuration 30
7.4.2 Experiment Execution 30
7.4.3 Host Monitoring 31

7.4.4 Result Storage and Export 31
7.5 Level 2: Bare-Metal Kubernetes Deployment 31
7.5.1 Kubernetes Deployment 32
7.5.2 TPC-DS Workload 32
7.5.3 Custom Dataset 33
7.6 Level 3: GKE Worker Scaling and Full Suite 33
7.6.1 GKE Deployment 33
7.6.2 Worker Scaling Study 34
7.6.3 Full Suite at Four Workers 35
7.6.4 TPC-H SF100 Query Performance 36
7.7 Cross-Environment Portability 36
7.8 Framework Overhead 37
7.9 Evaluation Limitations 38
7.10 Summary 38

8 Conclusion 39
8.1 Summary 39
8.2 Reflection 39
8.3 Extending TriBench 40
8.4 Future Work 40

Appendices 41

A CLI Command Reference 41

B Getting Started Guide 43

C Requirements Justification 45

D Result Analysis Command Output 46

Bibliography 52

1 Introduction

This chapter outlines the motivation for developing TriBench, a cross-platform benchmarking framework for SQL workloads on distributed data lakehouses. The key challenges in the problem space are identified, and the overall aims of the project are introduced.

1.1 Motivation

The data lakehouse has emerged as the dominant pattern for large-scale analytical data management, combining the low-cost storage of a data lake with the ACID guarantees and query performance of a data warehouse (Armbrust et al. 2021). Apache Trino (Sethi et al. 2019), a distributed SQL query engine originally developed at Meta, is widely used as the interactive query layer over such systems, with Apache Iceberg (Apache Software Foundation 2024*b*) providing the underlying open table format.

Benchmarking such a stack is considerably harder than benchmarking a traditional RDBMS. Rather than a single system under test, a Trino lakehouse deployment relies on several coordinated services that must all be configured, started, and stopped together. Each service carries hardware-dependent settings that must be tuned consistently before any experiment can run, and these settings must change when the deployment environment changes. Studying how the system scales, for instance by increasing the number of query processing nodes, requires changes across several components and a full redeployment of the stack. When this is done manually, the exact configuration used is rarely captured alongside the results, making it difficult for others to reproduce the same experiment on different hardware. An experiment tuned for a laptop may produce out-of-memory errors or skewed timings on a cluster node, with nothing in the results to indicate that the configuration was the cause.

This challenge is not unique to Trino. The Apache Spark ecosystem faced the same problem, and the PEEL framework (Boden et al. 2017) was developed to address it, automating the setup, execution, and teardown of multi-system experiments and packaging them into portable, reproducible units. No equivalent tool exists for Apache Trino. Without one, practitioners typically resort to ad-hoc scripts that embed machine-specific assumptions and store results as unstructured logs, making systematic comparison and reproduction difficult. TriBench is designed to fill this gap: a framework for defining, running, and reproducing SQL benchmark experiments across Trino deployments, from a local machine to a cloud cluster.

1.2 Aims

TriBench targets engineers and researchers who need to evaluate SQL query performance on distributed data lakehouses in a systematic and reproducible way. The framework is designed around five goals. First, it must run SQL benchmark workloads against a Trino cluster from a self-contained experiment definition, supporting both industry-standard suites and user-supplied queries. Second, it must collect hardware and system resource metrics concurrently with query execution, capturing data at the host level and, on Kubernetes, at the individual pod level. Third,

results must be stored in a structured form that supports post-hoc analysis and comparison across runs. Fourth, the same experiment definition must run unchanged across deployment environments by changing only a deployment configuration, covering both a local single-node setup and a multi-node cloud cluster. Fifth, all experiment artefacts must be packaged into a portable, self-contained unit that can be archived and reproduced on a different machine.

1.3 Outline

- **Chapter 2** introduces the foundational concepts: the data lakehouse architecture, Apache Iceberg, Apache Trino, and the principles of reproducible benchmarking.
- **Chapter 3** surveys related work, examining existing benchmarking frameworks, in particular PEEL for Apache Spark, and identifying the gap that TriBench addresses.
- **Chapter 4** presents a structured set of functional and non-functional requirements derived from the aims above.
- **Chapter 5** describes the design of TriBench: its five-layer architecture, configuration system, monitoring subsystem, and result storage schema.
- **Chapter 6** explains how the design was realised in Python, covering key libraries, the execution pipeline, and implementation challenges.
- **Chapter 7** evaluates TriBench across three deployment environments, assessing requirement coverage, performance, and reproducibility.
- **Chapter 8** concludes with a summary, critical reflection, limitations, and directions for future work.

2 Background

2.1 Data Lakehouse Architecture

The evolution of modern analytics platforms can be understood as a progression from data warehouses to data lakes, and finally to the data lakehouse pattern (Armbrust et al. 2021). Early data warehouses centralised curated data for decision support and business intelligence (Kimball and Ross 2013). They typically used a schema-on-write approach: data was transformed on ingest into a fixed model designed for predictable analytical queries. This provided strong governance and excellent performance, but often at the cost of tight coupling between storage and compute, expensive scaling as data volumes grew, and limited support for semi-structured or unstructured data.

Data lakes emerged to address these constraints by separating storage from compute. Raw data could be stored cheaply in generic file systems or object stores, commonly using open columnar formats (Abadi et al. 2008; 2013) such as Apache Parquet (Vohra 2016). This schema-on-read model improved agility and reduced cost, but it introduced new problems: without a transactional coordination layer and consistent metadata, concurrent writers can leave a dataset in an inconsistent state, and query engines may be forced into large scans because they lack reliable file- and partition-level statistics.

The *data lakehouse* approach aims to combine the openness and cost profile of data lakes with the management properties of warehouses. It does so by introducing an open table format layer that maintains transaction state and query-relevant metadata alongside the data files. Apache Iceberg (Apache Software Foundation 2024*b*) is a prominent implementation. An Iceberg table is defined by a metadata tree: a *metadata file* points to the current *snapshot*, which references a *manifest list* of *manifest files* describing the underlying Parquet objects, including column statistics. This metadata enables optimisations such as partition pruning and predicate pushdown, substantially reducing the amount of data that must be read for many analytical queries.

To make these tables discoverable via SQL, Iceberg deployments use a catalogue service that maps logical table names to metadata locations. In TriBench's stack, the Hive Metastore Service (HMS) (Thusoo et al. 2009) fulfils this role, persisting catalogue state in a relational backend. Apache Trino (Sethi et al. 2019), the distributed SQL engine that forms the compute layer of the stack, consults the catalogue during query planning to resolve the current metadata pointer for each referenced table. Trino operates under a coordinator–worker model: a coordinator node parses and plans queries, then distributes execution as parallel tasks across worker nodes that read Parquet data directly from object storage. This architecture allows the same experiment to scale from a single-node Docker deployment to a multi-node Kubernetes cluster by adjusting the worker replica count, without any change to the query workload or experiment specification.

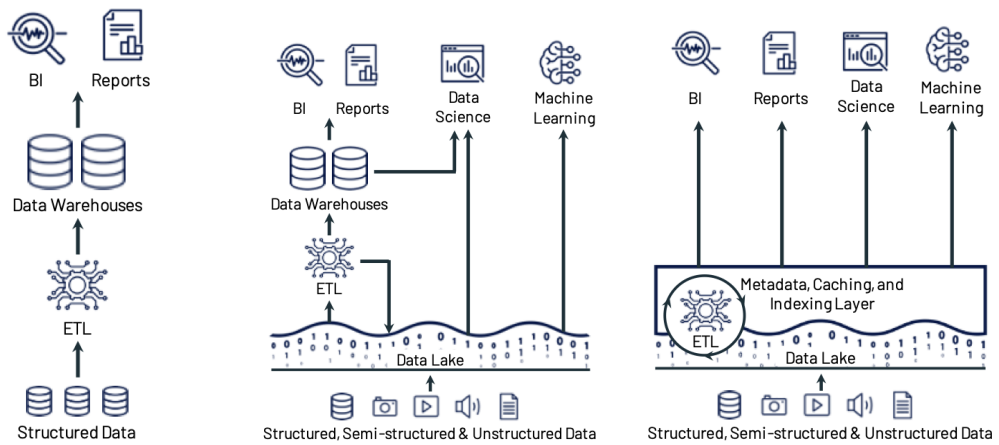

(a) Data Warehouse *(b) Data Lake* *(c) Data Lakehouse*

Figure 2.1: *Evolution of data platform architectures to today's two-tier model (a-b) and the new Lakehouse model (c). Armbrust et al. (2021)*

2.2 Benchmarking for Distributed SQL Systems

According to Gray (1993) and Huppler (2009), a useful benchmark must possess four properties: relevance, portability, scalability, and simplicity. The Transaction Processing Performance Council (TPC) provides the industry-standard suites that satisfy these criteria for analytical SQL workloads; the specific benchmarks used in this work are described in Section 2.3.

Benchmarking distributed SQL engines introduces significant complexity beyond single-node systems. Because queries execute as parallel jobs across multiple machines, end-to-end latency reflects not just computational work but also scheduling overhead, network exchanges, and remote storage throughput. This makes reproducibility (discussed in Section 2.4) a central concern: without tight control over environment, data layout, and execution parameters, results from one run cannot be meaningfully compared with another.

2.3 TPC Benchmark Standards

Founded in 1988, the Transaction Processing Performance Council (TPC) is a non-profit organisation dedicated to developing and maintaining data-centric benchmark standards and disseminating objective, verifiable performance data to industry (Transaction Processing Performance Council 2023). TPC benchmarks are divided into *Enterprise* benchmarks (technology-agnostic, specification-based standards where vendors choose their own implementation) and *Express* benchmarks, which are kit-based with shorter development cycles.

TPC-H (Transaction Processing Performance Council 2023) is a decision-support benchmark that models a supply-chain scenario with 22 business-oriented ad-hoc queries. The dataset and queries are designed for broad industry relevance, exercising decision-support systems that examine large volumes of data and execute queries with a high degree of complexity. Scale is controlled by a scale factor (SF), where SF 1 generates approximately 1 GB of raw data, SF 10 generates 10 GB, and so on, across eight related tables. TPC-H is the primary built-in benchmark suite shipped with TriBench: all 22 queries are included, pre-adapted to Trino-compatible SQL and parameterised by Iceberg catalogue and schema name so they can be used out of the box against a loaded dataset.

TPC-DS (Nambiar and Poess 2006; Poess et al. 2017; Transaction Processing Performance Council 2021) extends the decision-support model to a complex retail scenario with 99 queries testing advanced SQL features such as window functions, complex joins, and nested subqueries. TPC-DS is the second built-in benchmark suite: its query files are similarly stored in TriBench's directory, and TriBench's dataset registry and query-loader are structured to load and execute them against an Iceberg schema in exactly the same way as TPC-H, without any framework modification. Both suites form the primary workloads in the evaluation (Chapter 7).

2.4 Reproducibility in Systems Research

Jain (1991) defines reproducibility as the ability for an independent experimenter to obtain consistent results under identical conditions. In systems research, meeting this definition in practice requires capturing not just the software configuration but also the hardware environment, dataset state, and execution parameters, all of which interact in a distributed SQL system.

The distributed systems research community has long recognised this challenge (Collberg et al. 2016). A typical performance study requires coordinating multiple interdependent services (storage, catalogue, and compute engine), each with its own configuration. Without a systematic means of recording and replaying these configurations, reported results are difficult to verify and cannot be reproduced on a different machine. Boden et al. (2017) articulated this directly: practitioners are typically forced to choose between manually steering experiment lifecycles, an error-prone and non-reproducible process, or writing ad-hoc scripts that embed environment-specific assumptions. Both approaches produce artefacts that cannot be shared or re-run without significant manual effort.

Achieving reproducibility in distributed SQL benchmarking requires strict control over four dimensions that directly motivate the design of TriBench:

1. **Environment Specification:** Precise control over cluster configuration, including worker counts, CPU and memory limits, and JVM heap settings.
2. **Data Consistency:** Loading the dataset at an identical scale factor with the same physical layout and Iceberg partitioning strategy across runs.
3. **Execution Control:** Standardising the warmup policy, query execution order, and concurrency model so that variance between runs reflects only system behaviour.
4. **Measurement Capture:** Recording both query-level statistics (wall-clock time, CPU time, bytes scanned) and concurrent system-level telemetry (CPU utilisation, memory pressure). Capturing both is essential, as resource saturation often manifests in hardware metrics before it is visible in query latency.

The TriBench design maps each of these dimensions to a specific mechanism, as described in Chapter 5.

3 Related Work

This chapter surveys benchmarking frameworks and deployment technologies informing TriBench's design. Section 3.1 reviews specialised frameworks for distributed data processing: PEEL, LST-Bench, and DiPerF. Section 3.2 describes the Trino–Iceberg lakehouse stack and Kubernetes container orchestration.

3.1 Benchmarking Frameworks for Distributed Systems

3.1.1 PEEL

PEEL (Boden et al. 2017; Alexandrov et al. 2015) is a framework developed at TU Berlin for reproducible experiments on distributed data processing systems such as Apache Spark (Zaharia et al. 2016), Flink (Carbone et al. 2015), and Hadoop (Shvachko et al. 2010). It structures the benchmarking process into a continuous four-stage life cycle: Define, Execute, Analyse, and Share, as illustrated in Figure 3.1.

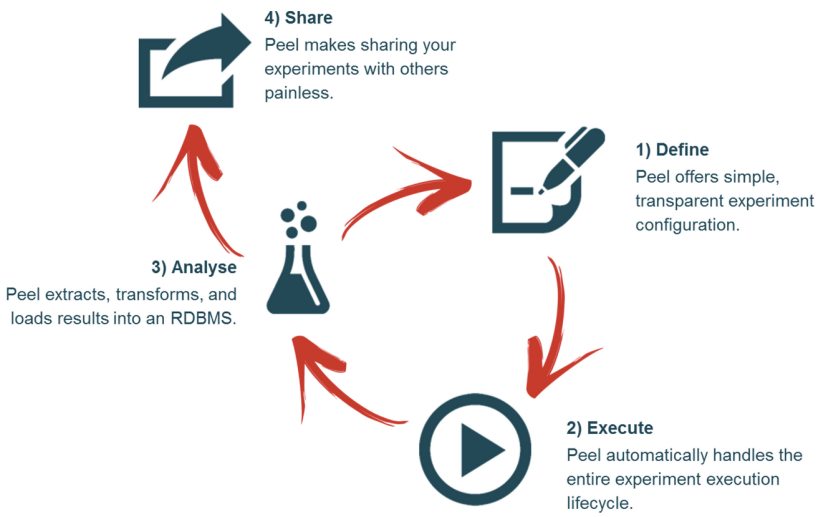

Figure 3.1: *The four stages of the PEEL experiment life cycle. Figure from Alexandrov et al. (2015).*

In the **Define** stage, experiments are configured through a layered HOCON (Lightbend 2023) hierarchy that separates framework defaults, bundle-level settings, and host-specific overrides. This allows the same experiment to run on different hardware without modification, a design pattern TriBench applies to the Trino lakehouse stack. In the **Execute** stage, PEEL automatically starts and stops each service based on a declared lifecycle scope, and triggers hardware telemetry collection alongside execution. In the **Analyse** stage, results are extracted from log files into a relational database, which feeds into the **Share** stage. The unit of sharing is the *bundle*: a

self-contained, portable directory packaging workload applications, configuration files, datasets, and result logs. Bundles can be pushed to remote clusters and shared with collaborators, who can reproduce the full experiment without additional setup.

PEEL's bundle paradigm and layered configuration model are proven design patterns, but two fundamental gaps prevent their application to modern lakehouse stacks. PEEL targets Hadoop-era engines on bare-metal or YARN, with no support for Kubernetes or the Trino–Iceberg–MinIO composition. Its telemetry is post-hoc, collected from log files after runs complete. TriBench applies these proven design patterns to solve a different problem: end-to-end lifecycle management of the complete Trino lakehouse stack across Docker and Kubernetes deployments, with concurrent hardware telemetry and per-query statistics captured from Trino's REST API during execution. No existing framework provides this capability for Apache Trino.

3.1.2 LST-Bench

LST-Bench (Camacho-Rodríguez et al. 2024) evaluates Log-Structured Tables (Delta Lake (Armbrust et al. 2020), Apache Iceberg, Apache Hudi (Apache Software Foundation 2024*a*)) on cloud object storage. Developed by Microsoft Research, it extends TPC-DS with composable workload patterns addressing LST-specific characteristics: longevity, resilience, read/write concurrency, and time travel. It quantifies degradation rate as delta files accumulate and automates telemetry from compute engines and cloud storage APIs to measure latency, throughput, storage efficiency, and compute utilization. LST-Bench compares table format performance, assuming engine and storage infrastructure are already provisioned. It does not manage deployment lifecycle of Trino, Hive Metastore, or object storage, nor provide bundle-based isolation or declarative profiles. TriBench is complementary: it provides the reproducible deployment harness and configuration management for executing LST-Bench workloads across environments.

3.1.3 DiPerf

DiPerF (Raicu et al. 2006) automates performance measurement of Grid computing networked services (GridFTP, GRAM, WS-MDS from Globus Toolkit). DiPerF coordinates distributed client machines concurrently accessing a target service, collects per-client response time and throughput, performs clock reconciliation across geographically dispersed nodes, and aggregates results into service-level statistics (fairness, saturation point). Its architecture cleanly separates concerns: the controller distributes client code to testers, testers time requests, and the offline analyser summarises performance samples, influencing TriBench's orchestration/monitoring/storage layer separation. However, DiPerF assumes the target service is deployed and running; it provides no provisioning or teardown mechanism. TriBench addresses this gap: it manages the full lakehouse service lifecycle (Trino, Hive Metastore, MinIO, PostgreSQL), captures per-query SQL statistics from Trino's REST endpoint alongside concurrent hardware telemetry, and uses bundle abstraction and HOCON hierarchy to guarantee reproducibility across deployments. DiPerF measures service response to distributed load; TriBench measures lakehouse end-to-end performance under controlled analytical workload.

3.2 Lakehouse Technologies and Container Orchestration

TriBench benchmarks the Trino–Iceberg lakehouse stack across Docker (Docker Inc. 2024) and Kubernetes deployments. This section surveys the core technologies under test and the container orchestration platform enabling reproducible distributed deployments.

3.2.1 Apache Iceberg

Apache Iceberg is an open table format that manages how data files are organised, versioned, and exposed as a single queryable table (Apache Software Foundation 2024*b*). To achieve this, Iceberg separates concerns into distinct catalog, metadata, and data layers, as illustrated in Figure 3.2. At the top, the Iceberg Catalog maintains a pointer to the current metadata file. This metadata layer records the schema, partitions, and per-file statistics. It branches down into manifest lists and individual manifest files, which ultimately track the actual data files residing in the data layer. Query engines traverse this hierarchical metadata to safely skip irrelevant files before reading the underlying data, drastically reducing unnecessary I/O.

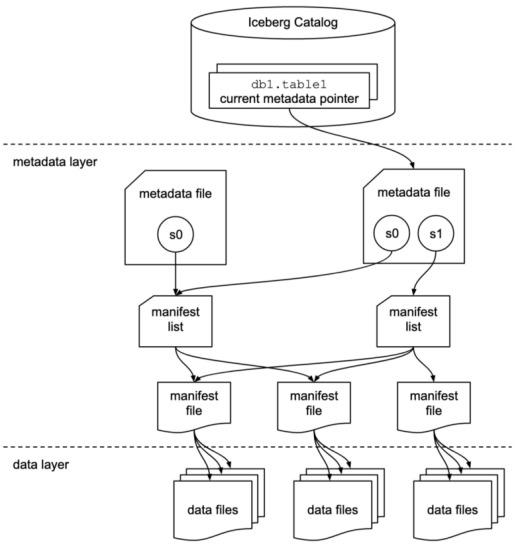

Figure 3.2: *Iceberg table format. Figure from Areca Data (2022).*

This layered architecture natively supports snapshot-based versioning, represented by the s0 and s1 states in the diagram's metadata files. Each write produces a new snapshot and atomically updates the catalog's metadata pointer, ensuring that readers always see a consistent view even during concurrent writes. For TriBench, Iceberg ensures that TPC-H and custom datasets are loaded in a consistent, versioned state, meaning benchmark results purely reflect query engine behaviour rather than underlying data inconsistencies.

3.2.2 Apache Trino

Apache Trino (formerly PrestoSQL) is a distributed SQL query engine (Sethi et al. 2019; Fuller et al. 2021) built on a strictly separated coordinator–worker architecture, as detailed in Figure 3.3. The lifecycle of a query begins when a Client submits a statement to the Coordinator. Inside the Coordinator, the query passes through a Parser & Analyzer, a Planner, and a Scheduler, which jointly translate the SQL into an optimised, parallelisable execution plan. The Scheduler then distributes these tasks to multiple Worker nodes. The workers perform the heavy lifting: they execute the plan, read data directly from the underlying storage, and stream results back. Because the architecture decouples query coordination from data processing, scaling the cluster simply means adding more workers; the query and experiment definitions stay unchanged.

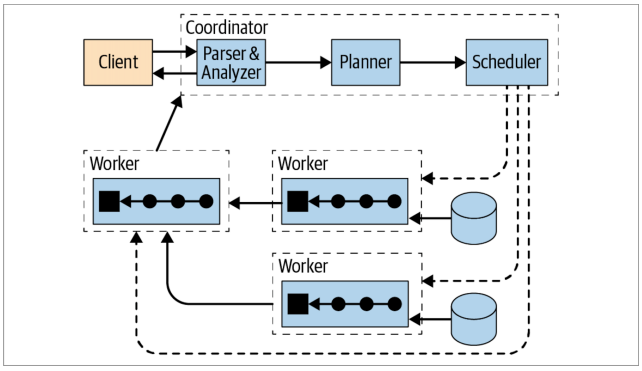

Figure 3.3: *Trino architecture overview with coordinators and workers. Figure from Trino Project (2021).*

As the workers execute these scheduled tasks, CPU, memory, disk I/O, and network resources heavily dictate Trino's overall performance. Workers consume CPU through sorting, hashing, and aggregation. Large joins and aggregations can exhaust a worker's available memory, causing disk spillage and sharply increasing query time. Furthermore, network throughput becomes a primary bottleneck when workers must exchange intermediate results between execution stages (indicated by the data flow arrows between workers in the diagram). These complex, distributed interactions make concurrent hardware monitoring essential, as a slowdown may originate in any of these interconnected resources. To provide visibility into this execution, Trino exposes per-query statistics (wall-clock time, CPU time, and bytes scanned) via a REST API endpoint after each query completes. TriBench collects these statistics and stores them alongside concurrent hardware metrics to precisely identify performance bottlenecks.

3.2.3 Kubernetes for Stateful Data Systems

Kubernetes (Burns et al. 2016) is the container orchestration platform enabling TriBench distributed benchmarks. The Docker backend runs all services on a single machine; the Kubernetes backend distributes Trino across cluster nodes: the coordinator pod handles query planning and orchestration, while worker pods execute data processing tasks. This allows TriBench to provision any number of Trino workers to a cluster by changing a single configuration parameter, enabling performance scaling measurement without modifying experiment definitions.

Kubernetes manages the full lakehouse stack as a coordinated unit (Trino coordinator, Trino workers, Hive Metastore, PostgreSQL catalogue, MinIO object storage), providing stable network identities, persistent storage, and service discovery so components locate and connect regardless of scheduling node. This abstraction allows identical experiment specifications to run unmodified on local kind clusters or multi-node production environments. Switching between clusters requires only changing the active configuration profile; TriBench sets the correct kubectl context automatically, leaving experiment definitions unchanged. Chapters 5 and 6 discuss TriBench's Kubernetes deployment strategy.

4 Analysis and Requirements

This chapter derives a prioritised set of functional and non-functional requirements from user scenarios. Requirements were elicited through supervisor discussions, a prototyping phase validating Kubernetes orchestration and concurrent monitoring, and lessons from PEEL (Section 3.1.1), LST-Bench (Section 3.1.2), and the gap identified in Section 1.1.

4.1 User Scenarios

Three scenarios model progressively more complex use cases, corresponding to the three evaluation levels in Chapter 7.

Scenario 1: Local baseline. A researcher wants to run TPC-H experiments against a Trino data lakehouse but has no prior experience configuring Iceberg, Hive Metastore, or MinIO. They create a bundle, activate the Docker profile, and run TPC-H queries. TriBench provisions the full lakehouse stack automatically, performs one warmup and three measured runs per query with host monitoring, then tears down. Results are exported to CSV for inspection.

Scenario 2: On-premises cluster study. A researcher benchmarks Trino on a self-managed Kubernetes cluster using a two-worker profile. They run queries on TPC-DS SF10 dataset and a custom dataset, with one warmup and five measured repetitions. TriBench collects per-pod CPU and memory alongside host telemetry.

Scenario 3: Cloud scaling study and sharing. An engineer runs a worker scaling study on GKE, executing TPC-H at one, two, and four workers to observe how query execution time changes with cluster size. They then stress test the framework via running a full suite including TPC-H at SF 100, TPC-DS at SF 10, and a custom workload runs unattended overnight, with the cluster fixed at four workers. The complete experiment is archived and shared; a colleague can reproduce it without any additional setup.

4.2 Requirements Set

Requirements are prioritised using MoSCoW (Clegg and Barker 1994): **M** must-have, **S** should-have, **C** could-have, **W** won't have.

4.2.1 Functional Requirements

Workload Execution (Goal 1). *Execute SQL workloads against Trino from a declarative definition, supporting standard suites and custom queries.*

1. **M1 – Declarative execution:** Experiments are orchestratable via YAML files without modifying framework code.
2. **M2 – Benchmark workloads:** TPC-H and TPC-DS are supplied; custom queries and user datasets are fully supported.
3. **M3 – Execution control:** Warmup runs, measured runs, timeouts, and validation thresholds are configurable.

4. **S1 – Suite abstraction:** Multiple experiments can be grouped into suites executable as a single CLI operation.

Hardware Monitoring (Goal 2). *Sample host CPU, memory, and I/O concurrently with query execution, alongside per-pod Kubernetes telemetry and query engine statistics.*

1. **M4 – Telemetry collection:** Host-level resources are sampled concurrently while per-query engine statistics are retrieved after each execution.
2. **S2 – Cloud metrics:** Kubernetes pod metrics are collected and normalised alongside host telemetry.

Result Storage (Goal 3). *Persist results and telemetry non-destructively, supporting cross-run analysis and export to common formats.*

1. **M5 – Structured storage:** Results and telemetry are stored non-destructively in a relational schema.
2. **M6 – Data export:** Stored runs can be exported to CSV, JSON, and Parquet.
3. **M7 – Built-in analysis:** Basic statistical evaluation including performance, query comparison, and monitoring summaries is provided.

Deployment Backends (Goal 4). *Run the same experiment unchanged across Docker and Kubernetes by switching only a host configuration profile.*

1. **M8 – Lifecycle abstraction:** Setup, execution, and teardown are automated across interchangeable Docker and Kubernetes backends.
2. **M9 – Configuration profiles:** Users switch deployment environments without modifying experiment definitions.
3. **W1 – Flexible stack:** Supporting alternative SQL engines is out of scope.

Bundle Abstraction (Goal 5). *Package all experiment artefacts into a self-contained bundle that can be archived and reproduced on any machine.*

1. **M10 – Self-contained bundles:** Benchmarks are packaged as shareable bundles containing all configurations and datasets.
2. **M11 – Bundle management:** The CLI scaffolds and activates bundles to direct all subsequent operations.

4.2.2 Non-Functional Requirements

1. **M12 – Reproducibility:** Identical experiments on identical hardware produce consistent results.
2. **M13 – Portability:** The framework executes unmodified across macOS and Linux.
3. **M14 – Observability:** Structured timestamped logs are emitted at configurable verbosity levels.
4. **S3 – Low overhead:** Background monitoring perturbs query execution times by less than 5% by default.
5. **C1 – Extensibility:** New deployment backends can be integrated via an abstract system interface.

4.3 Summary

This chapter presented three user scenarios and a prioritised requirements set covering workload execution (**M1–M3**, **S1**), hardware monitoring (**M4**, **S2**), result storage (**M5–M7**), deployment backends (**M8–M9**), bundle abstraction (**M10–M11**), and non-functional qualities of reproducibility, portability, observability, and low overhead. Chapter 5 presents the architecture satisfying these requirements; Chapter 7 provides evidence each was met, and a complete traceability table mapping every requirement to its evaluation evidence is provided in Appendix C.

5 Design

TriBench is organised as a five-layer architecture with two cross-cutting concerns. Each layer satisfies a distinct subset of requirements from Chapter 4; the design decisions around the configuration model, backend abstraction, and concurrent monitoring are motivated below alongside each component.

5.1 Architectural Overview

TriBench's five layers are shown in Figure 5.1: CLI, configuration, orchestration, systems, and data lakehouse. Control flows downward: a CLI command triggers configuration resolution, which parameterises orchestration, which drives system lifecycle management, which starts and stops the lakehouse services. Telemetry flows upward via two cross-cutting concerns: monitoring spans layers three to five, and result storage spans layers three to four.

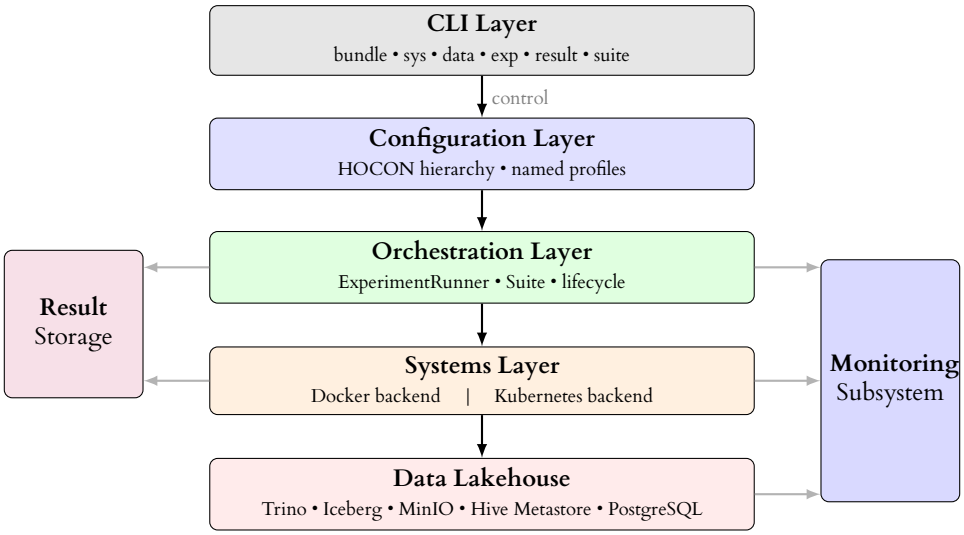

Figure 5.1: *Five-layer architecture. Control flows downward through the layers; monitoring and result storage span layers three to five and three to four respectively as cross-cutting subsystems.*

The layering addresses three non-functional requirements. **M13** (platform portability) is met by confining all platform-specific logic to the systems layer, so nothing above it contains environment assumptions. **C1** (extensibility) is a consequence: new deployment backends require only a concrete implementation of the System abstract base class in the systems layer, with no changes elsewhere. This mirrors the separation-of-concerns rationale that Raicu et al. (2006) applied to their controller–tester–analyser split (Section 3.1.3). Two cross-cutting subsystems span the orchestration and systems layers: monitoring (Section 5.7) collects hardware telemetry concurrently with query execution (**M4**, **S2**), and result storage (Section 5.10) persists run records to a structured database inside the bundle (**M5**, **M6**).

5.2 Open Lakehouse Stack

The framework evaluates a complete open lakehouse stack. **Trino** (Section 3.2.2) serves as the primary distributed SQL engine, chosen because unlike Spark and Flink, it previously lacked a dedicated benchmarking harness like PEEL. The stack is completed by **Apache Iceberg** (Section 3.2.1) for table storage, **MinIO** (MinIO Inc. 2024) for S3-compatible object storage, and a **Hive Metastore** (Thusoo et al. 2009) backed by **PostgreSQL** (The PostgreSQL Global Development Group 2024) for catalogue management. Depending on the scale of the experiment, this entire stack can be deployed locally via **Docker Compose** or distributed across a cluster using **Kubernetes**. Figure 5.2 shows how the four services compose.

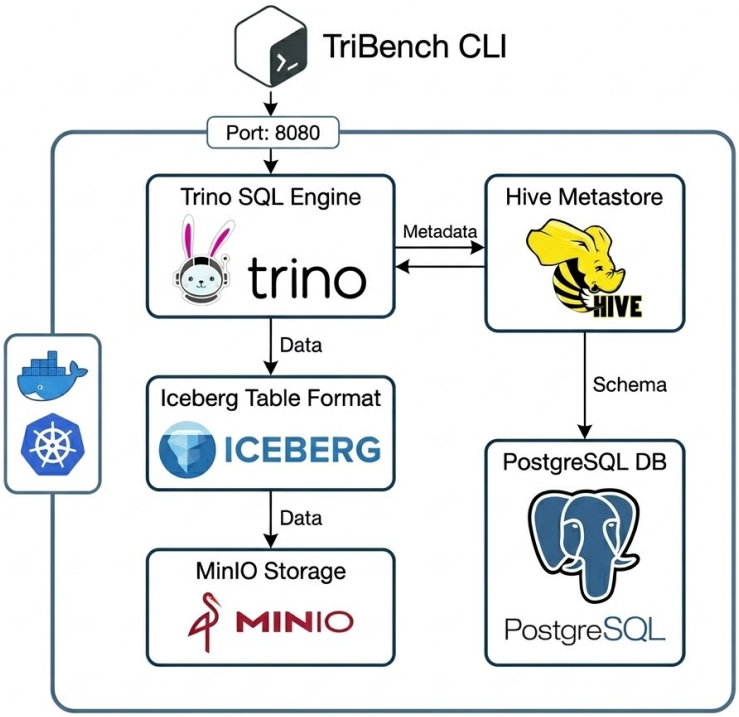

Figure 5.2: *TriBench open lakehouse stack. Trino queries tables registered in the Hive Metastore, whose state is persisted in PostgreSQL; data files are stored as Iceberg Parquet in MinIO. Logo sources: (Trino Project 2024; HPCwire 2025; Wikimedia Commons 2024; Open Policy Agent 2024; Icon Icons 2024).*

5.3 CLI Layer

The CLI layer is the single entry point for all user interaction with TriBench. It is built on Click (Pallets 2024*a*), which handles argument parsing, nested command groups, and context passing without global variables. The layer exposes seven command groups, each responsible for one area of the framework (Table 5.1). No framework code is imported or called outside these groups; all operations flow through CLI commands, ensuring consistent bundle resolution, logging, and error reporting regardless of which command is invoked.

Group	Responsibility	Requirements
`bundle`	Create, activate, validate, and archive bundles	**M10**, **M11**
`sys`	Start, stop, and tear down the lakehouse stack	**M8**
`data`	Load datasets and manage the dataset registry	**M2**
`exp`	Run experiments and select deployment profiles	**M1**, **M3**, **M9**
`result`	Export, display, and delete stored results	**M6**, **M7**
`config`	Switch profile; Inspects configurations	**M9**
`suite`	Run grouped experiment suites end-to-end	**S1**

Table 5.1: *TriBench CLI command groups and their requirement coverage.*

For a complete reference of all sub-commands and their arguments, see Appendix A.

5.4 Bundle Abstraction

A bundle packages together the configuration, experiments, datasets, queries, and results required for a particular benchmarking study. It is self-contained and can be archived and shared for reproducibility. Table 5.2 lists the top-level elements of a TriBench bundle.

Path	Overridable	Description
`bundle.yaml`	No	Bundle name and framework version
`config/`	Yes	HOCON configuration files
`experiments/`	Yes	Experiment definition files
`apps/`	Yes	SQL query files and workload assets
`datasets/`	Yes	Static datasets and the dataset registry
`log/`	No	Execution logs produced by the framework
`results/`	No	SQLite database accumulating all run results

Table 5.2: *Top-level layout of a TriBench bundle. Overridable paths can be remapped in bundle.yaml.*

The datasets directory additionally holds a registry file that maps logical dataset names to physical file locations, decoupling experiment definitions from paths. The results directory accumulates run records across invocations with each run assigned a unique identifier so repeated executions append without overwriting (**M5**). A typical end-to-end session follows this sequence of CLI commands; a full worked example is provided in Appendix B.

5.5 Configuration Layer

Each experiment in TriBench must be associated with a concrete set of configuration values covering both the deployment environment (Trino JVM heap, worker count, cluster address, backend type) and the workload (query timeout, run count, connection settings). Several problems arise with a naïve approach of configuring these manually per experiment:

Without a layered configuration model, machine-specific values (heap size, addresses, Kubernetes context) end up in experiment files; moving an experiment to a different machine requires editing those files, and related experiments that share most parameters duplicate values manually. TriBench adopts PEEL's layered model (Section 3.1.1), expressed in HOCON (Lightbend 2023), so a profile overrides only the parameters it needs without restating the full tree. Parameters are resolved through four ordered layers, each able to override the one below it (Figure 5.3).

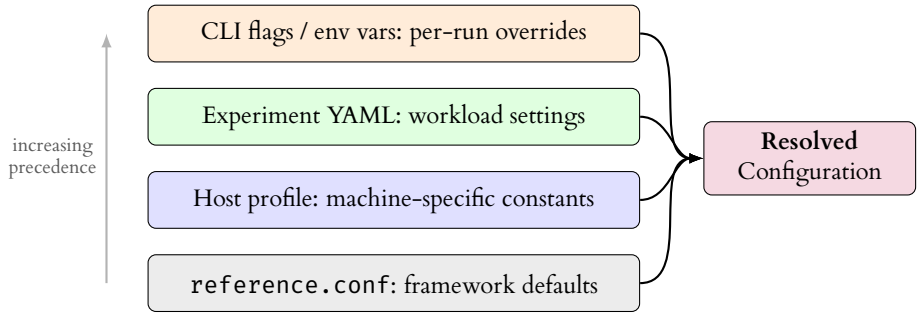

Figure 5.3: *Four-level configuration hierarchy (lowest to highest precedence). Each layer overrides keys from the one below; the resolved configuration is consumed by the orchestration layer.*

Layer 1 (reference.conf) is maintained by the framework and should not be edited. It declares every valid parameter once, so any key omitted at higher layers retains a safe default.

Layer 2 (named profile) is the recommended primary configuration point. An engineer creates one profile per deployment environment (a laptop Docker setup, a local kind cluster, a self-managed Kubernetes cluster) and stores it in `config/hosts/`. Each profile captures the machine-specific constants that should never appear in an experiment file: backend type, worker list, JVM heap, cluster address, and Kubernetes context. Profiles can live inside a bundle's own `config/hosts/` for bundle-local environments, or in the framework-level `config/hosts/` to be shared across bundles on the same machine.

Layer 3 (experiment YAML) is appropriate for parameters that vary between workloads within the same environment: run count, query timeout, or a Trino session property for one experiment. It must not contain machine-specific values; if it does, the experiment ceases to be portable across environments.

Layer 4 (CLI / environment variables) is intentionally narrow in scope. It suits one-off overrides and secret injection in CI/CD pipelines where credentials must not be stored in files. For normal benchmark use it should be avoided: a flag is not recorded in any file, making the run harder to reproduce from stored artefacts alone.

Named profiles at layer 2 are therefore the primary mechanism for deployment context switching and portability (**M9**, **M13**). A profile states only the keys it needs to override; all other parameters fall through to `reference.conf`. Switching profiles switches the entire deployment context (worker count, JVM heap, backend, cluster address) without touching any experiment file. The same experiment YAML runs unchanged on a single-node Docker stack or a multi-node Kubernetes cluster.

5.6 Orchestration Layer

5.6.1 Experiment Definitions

An experiment is declared as a YAML file inside the bundle's experiments directory. Each file specifies the workload independently of the deployment environment, so the same file runs unchanged across Docker, GPG, and GKE by switching only the active host profile. Listing 5.1 shows a representative experiment definition. Queries can be referenced as SQL files in the bundle's apps directory or written as inline SQL blocks directly in the YAML, which is used for custom or exploratory workloads.

```yaml
name: "tpch-sf1"
system: "trino"
connection:
  host: "localhost"
  port: 8080
  catalog: "iceberg"
  schema: "tpch_sf1"
runs: 3
warmup_runs: 1
timeout_seconds: 60
monitoring:
  enabled: true
  interval_seconds: 2.0
query_files:
  - "queries/tpch/queries/q01.sql"
  - "queries/tpch/queries/q06.sql"
validation:
  min_success_rate: 0.95
```

Listing 5.1: *Experiment definition for a TPC-H SF1 run.*

The key fields and their roles are summarised in Table 5.3. The workload fields (`runs`, `warmup_runs`, `timeout_seconds`) satisfy **M3**; the `query_files` and inline `queries` blocks satisfy **M1** and **M2**; the `monitoring` block activates the subsystem described in Section 5.7 (**M4**, **S2**).

Field	Purpose
`name` / `description`	Identifier written to the results database alongside each run
`system`	Backend type; currently `trino`
`connection`	Trino coordinator address, catalogue, and schema
`runs`	Number of measured repetitions per query
`warmup_runs`	Discarded warm-up executions before measurement begins
`timeout_seconds`	Per-query execution time limit
`query_files`	Paths to SQL files inside the bundle's apps directory
`queries`	Inline SQL blocks, used for custom workloads
`monitoring`	Enables background collection and sets sampling interval
`validation`	Sets minimum success rate to pass validation

Table 5.3: *Experiment definition fields and their roles.*

5.6.2 Execution Lifecycle

Before any experiment runs, its YAML file is merged with the active HOCON hierarchy (Section 5.5); the orchestration layer consumes the resolved result without knowing which layer contributed each value (**M1**, **M9**). It exposes two execution modes: a single experiment via `tribench exp run` and a suite via `tribench suite run`. Suite mode groups multiple experiments under a shared system lifecycle so the backend starts once and tears down after the final experiment, avoiding repeated cold-start latency (**S1**). In both modes the experiment lifecycle follows five phases in order, shown in Figure 5.4. The first and last of these are carried out by the systems layer; the orchestration layer contains no Docker- or Kubernetes-specific code, so switching backends requires only a profile change (**M13**).

Figure 5.4: Experiment execution lifecycle. Both single-experiment and suite modes follow these five phases in order.

5.7 Monitoring System

The monitoring system captures hardware telemetry concurrently with query execution without blocking the query thread. Each metric source implements a MetricCollector interface with start, stop, and collect methods; a central MonitoringSession coordinates all collectors and runs them in the background at the configured sampling interval. Three collectors are provided: a host resource collector that samples CPU, memory, disk I/O, and network I/O (**M4**); a Kubernetes pod collector that gathers per-pod CPU and memory and gracefully disables itself if the cluster Metrics Server is unavailable (**S2**); and a query statistics collector that reads Trino's per-query engine response after each execution (**M4**). Figure 5.5 shows the three-collector structure. The concurrent threading model and shutdown sequence are described in Chapter 6. A limitation of this polling approach is that resource spikes shorter than the sampling interval are not captured; for queries that complete in under one second, a single sample may represent the entire execution window.

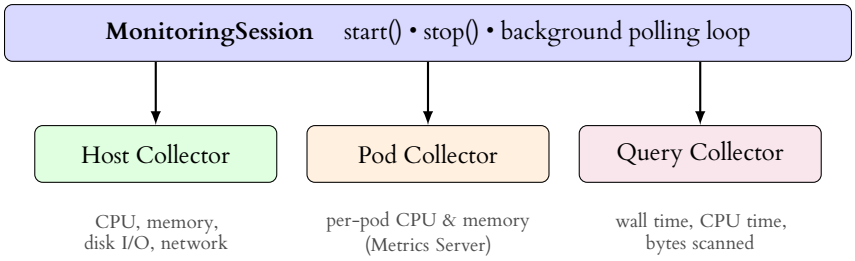

Figure 5.5: Monitoring subsystem. The MonitoringSession drives three collectors in a shared background polling loop; each implements the same MetricCollector interface.

5.8 Systems Layer

The systems layer manages the lifecycle of the infrastructure experiments run on. TriBench supports two interchangeable backends: Docker for single-machine deployment (**M8**) and Kubernetes for distributed clusters (**M8**, **M13**). The Kubernetes backend can target any available cluster, from a local KinD cluster for development to a managed service such as Google Kubernetes Engine, by pointing at the appropriate context in the active profile.

Both backends deploy the same four services: Trino, Hive Metastore, MinIO, and PostgreSQL. HMS, backed by PostgreSQL, provides the Iceberg catalogue layer; it was chosen over the Iceberg REST catalogue because it is the most widely deployed catalogue in self-managed Trino installations, so the stack reflects a common production configuration rather than a framework-specific one. MinIO acts as an S3-compatible object store, and because Trino's Iceberg connector uses S3 semantics, it provides identical connectivity to Amazon S3 without requiring cloud credentials or network access. This is what enables local reproducibility (**M12**): the full stack runs on a single machine with no external dependencies.

Both backends expose the same lifecycle operations to the orchestration layer: set up, start, stop, and tear down. The System abstract base class defines exactly these four methods plus a status

query; the orchestration layer calls only these and contains no Docker or Kubernetes API code of its own. Adding a third backend requires only a new subclass with no changes elsewhere (**C1**). Figure 5.6 shows the lifecycle state machine both backends implement.

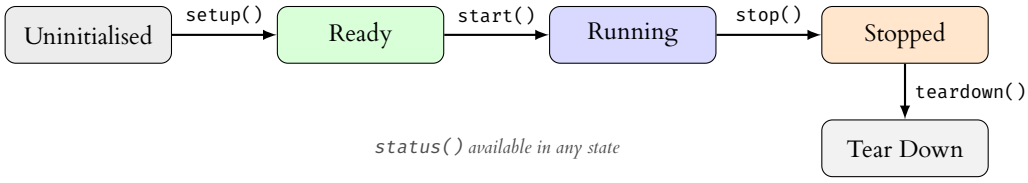

Figure 5.6: *System lifecycle state machine. Both backends implement the same four transitions.*

5.9 Dataset Loading

Before an experiment can run, benchmark data must be available as Iceberg tables in the lakehouse. TriBench provides two loading paths selected automatically based on the dataset type (**M2**). For TPC-H and TPC-DS, data is generated on demand without requiring any local files. For externally supplied datasets, a general-purpose pipeline accepts Parquet files, uploads them to object storage, and loads them into Iceberg through a temporary staging step. Both paths produce the same result: Iceberg tables registered in the Hive Metastore and ready to query. After loading, a provenance record is stored capturing the dataset version, tying each benchmark result back to the exact data used and supporting reproducibility across runs. The implementation of both paths is described in Section 6.5.

5.10 Result Storage

The results of every experiment run are stored in a relational database that persists across sessions. Each run produces three categories of records: per-query execution results containing timing and Trino engine statistics, a time-series of hardware telemetry samples captured during execution, and a run-level summary with overall status and resource totals. Runs are linked to a named experiment record, so executing the same workload twice appends a new run rather than overwriting the previous one.

The schema separates the experiment definition from its runs, so the same named workload can accumulate results across multiple invocations without duplication. Each run then owns two independent record types — per-query execution results and time-series monitoring samples — keeping query performance data and hardware telemetry linked but separately queryable. The stored data can then be explored using the built-in analysis commands or exported for use with external tools.

SQLite (Hipp 2024) is the default database for local use: it is embedded and serverless, so the bundle requires no external database process and results are stored inside the bundle directory with no configuration. PostgreSQL (The PostgreSQL Global Development Group 2024) is supported for shared or high-volume deployments; because the storage layer uses an ORM, switching between the two requires only a configuration change. Results can then be explored via built-in `tribench res analyze` subcommands covering statistics, performance comparison, regression detection, and scalability analysis, or exported to CSV, JSON, or Parquet for external tools (**M6**, **M7**). The implementation of these commands is described in Chapter 6.

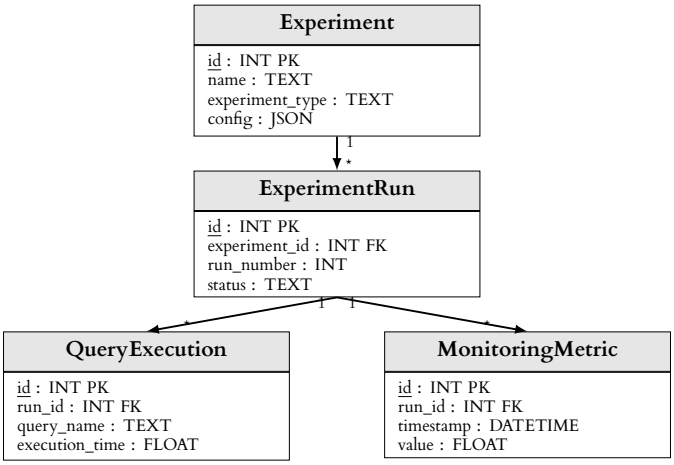

Figure 5.7: *Result storage schema (key fields only).*

5.11 Reproducibility Mapping

Section 2.4 identified four dimensions required for reproducible distributed SQL benchmarking. Named host profiles (Section 5.5) satisfy **environment specification** (**M8**, **M9**): switching profiles changes the full deployment context without touching experiment files. **Data consistency** (**M2**) is addressed by recording Iceberg snapshot identifiers and row counts after each dataset load (Section 5.10), tying every result to the exact data version used. **Execution control** (**M3**) is enforced by the experiment YAML, which fixes the warmup policy, run count, query order, and timeout (Section 5.6.1). **Measurement capture** (**M4**, **S2**) is handled by the background monitoring thread, which samples host and per-pod metrics concurrently with query execution and reads Trino engine statistics per query (Section 5.7).

5.12 Summary

This chapter covered TriBench's five-layer architecture, tracing execution from the CLI through configuration resolution and orchestration to the lakehouse services, with monitoring spanning the orchestration, systems, and data lakehouse layers, and result storage spanning the orchestration and systems layers. The next chapter details how each of these components was built in Python.

6 Implementation

Development followed an incremental strategy: the Docker backend and core execution loop were implemented first to establish a working end-to-end path, followed by the monitoring subsystem, the Kubernetes backend, and finally the bundle management system.

6.1 Development Approach

6.1.1 Language and Tooling

TriBench is implemented in Python 3.11. Python was chosen for a benchmarking framework targeting Trino: the official Trino client library, the Docker SDK, the Kubernetes client, and the data science ecosystem (SQLAlchemy, pandas, psutil, matplotlib) are all mature, first-class Python packages. The framework is distributed as a pip-installable package with a tribench console entry point, so a single pip install makes the CLI available globally in the active environment.

Dependencies are managed using a conda environment defined in environment.yml, chosen over a plain virtual environment because several dependencies, particularly PyArrow, require pre-compiled native binaries that conda resolves consistently across platforms. Key dependencies are Click 8 (Pallets 2024*a*), pyhocon, the official Trino Python client, psutil (Rodola 2024), the Docker SDK, SQLAlchemy 2.0 (Bayer 2024), Jinja2 (Pallets 2024*b*), pandas (McKinney 2010), and PyArrow.

6.1.2 Version Control and Testing

Development used Git with feature branches for each major subsystem: the Docker backend, the monitoring subsystem, the Kubernetes backend, and bundle management were each developed in isolation before being merged into main. Testing used pytest (pytest Contributors 2024) for unit tests covering configuration merging, bundle path resolution, and result storage. Integration testing for the Kubernetes backend used KinD (Kubernetes SIGs 2024) (Kubernetes in Docker), which runs a full multi-node cluster inside Docker containers on a single machine. KinD was chosen because it exercises the full kubectl, readiness polling, and port-forward pipeline locally without cloud credentials. A CI pipeline was not set up because the test suite depends on a running Docker daemon and the KinD cluster, neither of which are available on standard CI runners without privileged access.

6.2 CLI Implementation

The CLI layer is built on Click and exposes seven command groups as described in Section 5.3. Each command has access to a shared context object carrying three fields set once at startup: the active bundle path, a verbosity flag, and a dry-run flag. All subcommands read from this context rather than managing their own state, ensuring consistent bundle resolution and logging across every command. Bundle resolution follows a fixed priority order: the `-bundle` flag takes precedence, then a state file recording the last activated bundle, then upward directory traversal

looking for a bundle.yaml file, so commands work from inside a bundle directory without an explicit path.

6.3 Bundle Implementation

The bundle system centres on a manifest class that exposes the bundle's directory layout as resolved path properties. Every layer (CLI, configuration, orchestration, and result storage) calls these properties rather than constructing paths directly. This means bundles can be moved or renamed without breaking any part of the framework, as path resolution is centralised in one place. The active bundle is tracked by a state file named .tribench-bundle stored at the framework root. Locating this file follows the same upward-directory walk used by Git to find its .git directory, so commands can be run from any subdirectory within a bundle and the correct bundle is resolved automatically.

When a new bundle is created, the framework scaffolds the full directory structure and writes annotated template files: the config directory receives a commented reference profile, the experiments directory a minimal YAML template with every field documented inline, and the apps directory a placeholder queries subdirectory. Once a study is complete, bundle archive packages the bundle into a compressed tarball with the framework version pinned, so a recipient can re-run the full experiment without additional setup (**M11**).

6.4 Configuration Loading

Experiment configuration is loaded into an ExperimentConfig dataclass via a class method that applies four layers of precedence: the framework defaults in reference.conf form the base, the active host profile overrides deployment-specific parameters, the experiment YAML applies workload-specific settings, and CLI flags override last. Nested keys are merged recursively via a deep merge helper so that a profile overriding only trino.coordinator.jvm.heap does not erase the sibling keys that reference.conf provides for the same object. The result is a flat, fully-resolved object that the orchestration layer consumes without knowing which layer contributed each value.

The GCP GKE profile, for example, overrides only backend type, JVM heap, memory limits, and cluster coordinates, roughly a dozen lines against the 304 lines declared in reference.conf:

```
tribench {
  defaults.backend = "kubernetes"
  systems.trino {
    coordinator.jvm.heap = "8G"
    query.max_memory = "20GB"
    query.max_memory_per_node = "5GB"
    jvm_options = ["-server", "-XX:MaxRAMPercentage=80"]
  }
  workers = [{name = "worker-1"}, {name = "worker-2"}]
  kubernetes {
    context = "gke_tribench_us-central1-a_tribench-cluster"
    namespace = "tribench"
  }
}
```

Listing 6.1: *GCP GKE host profile, illustrating the sparse-override pattern.*

Two diagnostic commands aid debugging: tribench config show prints the fully merged parameter tree, and tribench config trace identifies which layer last set each parameter (**M14**).

6.5 Dataset Loading

Before experiments can run, benchmark data must be materialised into Iceberg tables accessible to Trino. TriBench provides two complementary loading paths, selected automatically based on the dataset type: a built-in path for TPC-H and TPC-DS that requires no local files, and a universal path for externally provided Parquet datasets. TPC-H issues a direct CTAS from Trino's built-in tpch connector, bypassing external storage entirely. TPC-DS generates data via DuckDB and writes Parquet to MinIO (Phase 1); external Parquet datasets are uploaded the same way. Both then stage through a temporary Hive external table (Phase 2) before a final CTAS into Iceberg (Phase 3). All paths conclude with a `$snapshots` provenance record. Figure 6.1 illustrates the three paths.

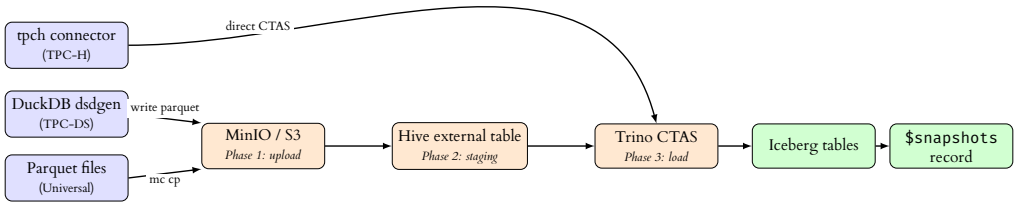

Figure 6.1: *The three dataset loading paths in TriBench.*

6.5.1 Built-in Loading

Trino ships with a built-in TPC-H connector that generates benchmark data on demand at any scale factor without requiring data to be physically written to disk first. TriBench exploits this by issuing a single CTAS statement from the virtual tpch connector schema directly into the Iceberg catalog, so the data is generated, streamed through the Trino coordinator, and written to MinIO as Iceberg Parquet files in one pass. Large fact tables such as lineitem and orders receive date-based partitioning automatically, which the evaluation experiments rely on for partition pruning tests.

TPC-DS has no equivalent built-in connector in Trino. Instead, TriBench uses DuckDB (Raasveldt and Mühleisen 2019), an embeddable in-process analytical database, as a generation engine: it calls the dsdgen routine at the requested scale factor, then writes each of the 24 TPC-DS tables as Parquet directly to MinIO over S3. Once all files are uploaded, Trino takes over and loads each table into Iceberg via the same CTAS-via-Hive-staging pattern described in Section 6.5.2. No local disk space is required beyond the DuckDB process memory, making this path suitable for containerised environments.

6.5.2 Universal Loader

For externally supplied Parquet datasets, TriBench provides a benchmark-agnostic three-phase loading pipeline implemented in the UniversalIcebergLoader class. The same code handles any dataset without modification, satisfying requirement **M2**.

Phase 1 – Parquet upload. Each Parquet file is copied to MinIO under a schema-namespaced prefix using the MinIO client. The loader checks whether the file is already present before uploading, so re-running a partially failed load is safe.

Phase 2 – Hive external table staging. Trino cannot read raw Parquet files directly from S3 without a table definition. The loader creates a temporary external table in the Hive Metastore

pointing at the uploaded file. Column definitions are derived automatically by reading the PyArrow schema embedded in the Parquet file footer and mapping each Arrow type to its Trino SQL equivalent. The staging table receives a millisecond timestamp suffix to prevent conflicts when two loads run concurrently.

Phase 3 – CTAS into Iceberg. The loader issues a CTAS statement that selects all rows from the Hive staging table and writes them into the Iceberg catalog. Optional partition columns, passed as a per-table mapping, are translated into an Iceberg partitioning array in the WITH clause. The staging table is dropped unconditionally in a finally block after each table completes.

```
CREATE TABLE iceberg.benchmark.lineitem
WITH (
  format = 'PARQUET',
  partitioning = ARRAY['month(l_shipdate)'],
  location = 's3://warehouse/benchmark/lineitem/'
)
AS SELECT * FROM hive.staging.lineitem_1710000000000;
```

Listing 6.2: *CTAS pattern used by the universal loader to copy data from Hive staging into an Iceberg table.*

After all tables are loaded, the loader queries the Iceberg `$snapshots` metadata table for each table and records the resulting snapshot identifier and commit timestamp alongside the row count. This lightweight provenance record ties each benchmark result back to the exact data version that produced it, supporting reproducibility across runs.

6.6 Experiment Execution

6.6.1 Single Experiment

The TrinoExperiment class implements the five-phase lifecycle described in Section 5.6.2. Within the execute queries phase, each query first runs for a configurable number of warmup iterations whose output is discarded, ensuring that JVM class loading and Iceberg metadata caching do not bias the measured results (Georges et al. 2007). Measured runs then proceed sequentially using the Trino Python client's cursor interface.

After each query, engine statistics are read from the cursor's stats attribute before the next execute call, as each call overwrites the previous stats. Listing 6.3 shows the relevant extraction:

```
cursor.execute(query)
rows = cursor.fetchall()
# Stats must be read before the next execute() call
stats = cursor.stats or {}
result = QueryResult(
   wall_time_ms = stats.get("elapsedTimeMillis"),
   cpu_time_ms = stats.get("cpuTimeMillis"),
   bytes_scanned = stats.get("processedBytes"),
   peak_memory = stats.get("peakMemoryBytes"),
   spilled_bytes = stats.get("spilledBytes"),
)
```

Listing 6.3: *Extracting Trino engine statistics from the cursor immediately after fetchall.*

6.6.2 Suite Execution

When running a suite, the ExperimentSuite class groups multiple experiments under a shared system lifecycle, passing the already-running system reference to each TrinoExperiment so it connects without triggering another start (**S1**). If any experiment fails, the suite logs it, continues with the remaining experiments, and reports a partial success rather than aborting.

6.6.3 Parallel Query Execution

For workloads that issue multiple queries concurrently, TriBench provides a parallel execution mode built on a thread-safe connection pool. The pool is backed by a `queue.Queue` and supports both eager pre-creation and lazy creation, configurable per experiment. Each worker thread in a `ThreadPoolExecutor` acquires a connection from the pool, executes its assigned query, releases the connection back to the pool, and appends its result to a shared list under a `threading.Lock`. Results are collected as they complete using `as_completed`, so a slow query on one thread does not delay the recording of fast queries on others. This design keeps the connection count bounded by the pool size regardless of the degree of parallelism, preventing connection exhaustion on the Trino coordinator.

6.7 Concurrent Monitoring

6.7.1 Threading Model

The monitoring subsystem runs in a background thread so resource sampling never blocks query execution. A MonitoringSession manages all collectors via a shared stop flag; the main thread spawns the background collector at session start, then sets the stop flag after all queries complete and calls `thread.join()` to wait for the final loop iteration to finish before flushing buffered samples to the database. Figure 6.2 shows the two-thread interaction.

```
def _collection_loop(self):
   while not self._stop_event.is_set():
      with self._metrics_lock:
         self._collected_metrics.extend(self.collect_metrics())
      self._stop_event.wait(self.config.interval_seconds)
```

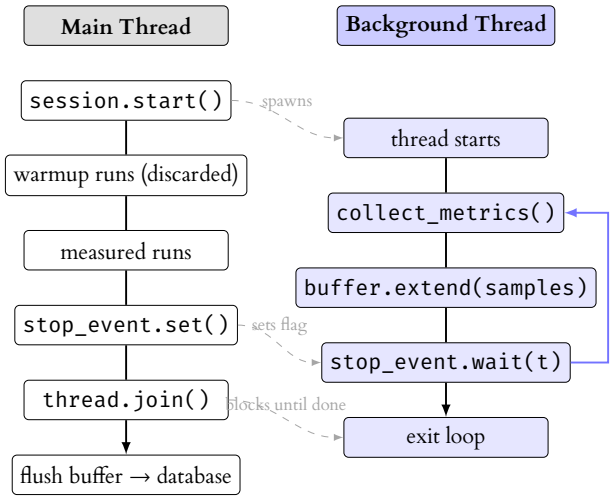

Figure 6.2: Monitoring thread lifecycle.

6.7.2 Host Resource Collector

The host resource collector uses psutil to sample CPU utilisation, memory, disk I/O, and network I/O at each interval. A key implementation detail is that psutil reports disk and network counters as cumulative totals since system boot. To produce per-experiment deltas, the collector records baseline counter values at session start and subtracts them from each subsequent reading:

$$\Delta_{\text{disk_read}} = C_{\text{now}} - C_{\text{baseline}}$$

Without this baseline subtraction, the reported figures would reflect the lifetime of the host process rather than the experiment duration. CPU utilisation does not require baseline subtraction because psutil already computes a percentage over the elapsed interval.

6.7.3 Kubernetes Pod Collector

On Kubernetes deployments, a second collector queries the cluster Metrics Server for per-pod CPU and memory at each interval. The collector issues a kubectl top pods command via subprocess and parses the output into normalised per-pod records, converting millicores and mebibytes into consistent numeric units. Subprocess was preferred over the Kubernetes Python metrics client because kubectl top is available on any cluster where the Metrics Server is deployed and avoids an additional client dependency. If the Metrics Server is unavailable, for instance because it was not deployed to the cluster, the collector catches the error at startup and disables itself gracefully, allowing the experiment to continue with host-level metrics only. This is the mechanism behind the GPG pod metric absence discussed in Section 7.5, where the cluster context was initially misconfigured.

6.8 Deployment Backends

6.8.1 Docker Backend

The Docker backend generates Trino configuration files from Jinja2 templates using the resolved profile; Jinja2 was chosen because its variable substitution maps directly to the resolved HOCON profile keys, avoiding a custom rendering step. Changing the JVM heap and restarting immediately produces an updated coordinator config. Once the containers are started, the backend performs a two-stage readiness check: the first stage polls the Docker SDK until each container reports running; the second polls the Trino coordinator's `/v1/info` REST endpoint and checks that the `starting` field has returned false, confirming the JVM is fully initialised before queries are submitted.

6.8.2 Kubernetes Backend

The Kubernetes backend follows the same Jinja2 template approach, generating service manifests from the active profile and applying them to the cluster via kubectl. Services are started in dependency order (MinIO first, then PostgreSQL, then Hive Metastore, then Trino), with each stage waiting for the previous to reach a ready state before proceeding. This ordering was important in practice: applying all manifests simultaneously caused Hive Metastore to fail on startup because the PostgreSQL catalogue it depends on was not yet accepting connections.

Once the cluster is running, the backend establishes a port-forward tunnel from a fixed local port to the Trino coordinator service. The tunnel runs as a background subprocess managed through a PID file, so it persists independently of the terminal session that started it. If the tunnel process dies, for instance due to a network drop, it can be restarted independently with tribench sys

port-forward without touching the cluster. However, if the tunnel dies silently mid-experiment, the orchestration layer reports a Trino connection error rather than identifying the tunnel as the cause, requiring manual diagnosis. The orchestration layer connects to the same local address regardless of whether the backend is Docker or Kubernetes, which is the concrete mechanism behind deployment portability (**M13**).

6.9 Result Storage

6.9.1 Persistence

Results are persisted using SQLAlchemy's ORM, which abstracts over both SQLite and PostgreSQL. Each invocation appends a new run record rather than overwriting existing data, so repeated runs accumulate and earlier results are never lost (**M5**). Each run is assigned a unique identifier combining the experiment name, a UTC timestamp, and a short random hex suffix to prevent collisions.

6.9.2 Export

Results can be exported in three formats via tribench res export. CSV produces two files per run: one for query execution records and one for monitoring metric samples, with a shared timestamp in the filename to keep them paired. JSON packages all records alongside the full resolved configuration snapshot into a single document, making it self-contained for archiving. Parquet targets downstream analysis with tools such as pandas or Apache Spark, where the columnar format allows efficient filtering over large numbers of runs (**M6**, **M7**).

6.9.3 Analysis Commands

The built-in analysis commands operate directly against the database without requiring an export step, making exploratory analysis faster during active experimentation. Five subcommands are provided under tribench res analyze:

statistics computes descriptive statistics across all measured runs for each query: mean, median, standard deviation, and P50/P95/P99 percentiles. IQR-based outlier detection (Jain 1991) flags individual runs that deviate significantly from the rest, which can indicate OS interference or resource contention during that run.

performance reports overall throughput in queries per second and total execution time for a single experiment run, with a per-query breakdown showing mean, min, and max latency. This gives a quick summary of the workload's execution profile without needing to inspect individual query results.

compare performs a head-to-head comparison between two experiment runs, for example a Docker run against a Kubernetes run, or SF 1 against SF 10. A two-sample t-test (Georges et al. 2007) determines whether observed latency differences are statistically significant at a configurable significance level; it was chosen because query execution times across repeated runs approximate a normal distribution, and small sample sizes (typically 3–5 runs) make non-parametric alternatives less powerful. This prevents visual noise in bar charts from being misinterpreted as meaningful differences.

regression identifies queries that have slowed beyond a configurable threshold (default 5%) with statistical significance and classifies each regression by severity: minor under 10%, moderate under 25%, major under 50%, and critical above 50%. A regression is only flagged when the percentage slowdown exceeds the threshold *and* a t-test confirms the difference is statistically significant, preventing noise in short run series from generating false alerts. The command exits

with a non-zero status code when regressions are detected, making it suitable for integration into a CI pipeline to catch performance changes automatically.

scalability compares two runs at different worker counts and computes speed-up ($T_{\text{baseline}}/T_{\text{scaled}}$) and parallel efficiency ($(S_{\text{actual}}/S_{\text{ideal}}) \times 100\%$) (Amdahl 1967) per query. An efficiency above 90% is classified as near-linear scaling; results below 50% point to coordinator bottlenecks or data skew.

Example output from all five subcommands is shown in Appendix D.

6.10 Summary

This chapter covered the key implementation decisions across each component: the four-layer HOCON merge, cursor-attached stats extraction, the two-stage Docker readiness check, dependency-ordered Kubernetes startup, and the five built-in analysis commands. Together these satisfy all Must-Have and Should-Have requirements, as confirmed in Chapter 7.

7 Evaluation

This chapter evaluates TriBench across five parts: unit testing, experimental setup, framework overhead, three end-to-end deployment scenarios, and a requirements coverage review.

7.1 Unit Testing

A suite of 267 unit tests was written using pytest pytest Contributors (2024), covering configuration, storage, query execution, statistical analysis, monitoring, CLI utilities, and bundle management. All external dependencies (Trino, Docker, Kubernetes, and PostgreSQL) are replaced with stubs so that tests run without any deployed infrastructure, with the result database held in memory. All 267 tests pass.

Table 7.1 shows line coverage per module across the tested components, giving 85% overall, above the 70% threshold recommended by Microsoft Microsoft (2024). Five modules reach 100%. The lower figures for query execution and monitoring reflect error-handling paths that only trigger under live network failures or when a background thread is forcibly interrupted; these are exercised by the end-to-end experiments in the sections that follow.

```
267 passed, 4 warnings in 1.02s
Name                                                    Stmts   Miss  Cover
-------------------------------------------------------------------------------
lib/tribench/analysis/statistical.py                       98     14    86%
lib/tribench/bundle/manifest.py                           170     22    87%
lib/tribench/cli/base.py                                  123      5    96%
lib/tribench/config/connection.py                         143     20    86%
lib/tribench/core/experiment.py                           102     15    85%
lib/tribench/core/result.py                                27      0   100%
lib/tribench/defaults.py                                  140     16    89%
lib/tribench/experiments/query_executor.py                192     68    65%
69
lib/tribench/monitoring/base.py                           161     20    88%
lib/tribench/storage/connection.py                         99     31    69%
lib/tribench/storage/models.py                            123      0   100%
lib/tribench/storage/result/__init__.py                     6      0   100%
lib/tribench/storage/result/experiment_store.py            48      7    85%
lib/tribench/storage/result/metric_store.py                72      7    90%
lib/tribench/storage/result/query_store.py                 21      2    90%
lib/tribench/storage/result/result_storage.py              42      2    95%
lib/tribench/storage/result/run_store.py                   59     10    83%
-------------------------------------------------------------------------------
TOTAL                                                    1626    239    85%
```

Figure 7.1: Test coverage report from `pytest -cov=. -cov-report=term`

7.2 Evaluation Methodology

Rather than a user study or comparative evaluation, this chapter takes a demonstration-based approach: each experiment shows that TriBench operates correctly and meets its requirements under progressively complex conditions. A user study was not conducted because meaningful feedback would require evaluators with direct experience running distributed SQL workloads on Trino, a sufficiently specific profile that recruiting an adequate sample was not practical within the scope of this project.

The evaluation is therefore structured as three progressively complex deployment levels: a local Docker stack, the University of Glasgow GPG Kubernetes cluster, and a managed GKE cluster. Each level introduces additional scale and environmental complexity, following the staged evaluation strategy used by Raicu et al. (2006) to isolate performance characteristics at each tier before moving to the next. Workloads are drawn from TPC-H and TPC-DS, chosen because they are industry-standard, vendor-neutral benchmarks with well-defined correctness criteria and reproducible data generation (Huppler 2009). Each query is preceded by a configurable number of warmup runs whose results are discarded, eliminating JVM class-loading and metadata caching effects from the measurements (Georges et al. 2007). Measured runs are then repeated and their timings recorded independently, so that variability across runs is visible rather than averaged away. Following the three deployment levels, a cross-environment portability study runs identical experiment definitions across all three platforms to confirm that only the host profile changes between environments, and a dedicated overhead experiment quantifies the performance impact of the background monitoring subsystem.

7.3 Experimental Setup

7.3.1 Local Environment (Level 1)

Level 1 experiments run on an Apple MacBook Pro with an M4 Pro chip, 24 GB unified memory, and an NVMe SSD running macOS Sequoia. Docker Desktop provides the container runtime, with Trino 434 as the query engine and a 4 GB coordinator heap. The TPC-H scale factor 1 dataset is stored as Parquet files in MinIO and registered through Hive Metastore, backed by PostgreSQL. The entire stack is brought up with a single command; step-by-step setup instructions are provided in Appendix B. All experiment definitions, configuration, datasets, and results are packaged into a single bundle for portability and reproduction.

7.3.2 University of Glasgow Kubernetes Environment (Level 2)

Level 2 experiments run on the University of Glasgow's GPG Kubernetes cluster. Each physical node runs a single Kubernetes node, so each Trino worker pod has exclusive access to a dedicated machine with no resource sharing. Each node has dual Intel Xeon E5-2640 processors (16 cores at 2 GHz), 64 GB RAM, 300 GB local disk, and a 10 Gb Ethernet interconnect. TriBench connects to the cluster automatically using port forwarding, and the Kubernetes Metrics Server is deployed to collect per-pod CPU and memory usage. A two-worker configuration is used with a 16 GB coordinator heap. The TPC-DS scale factor 10 dataset is loaded alongside a custom domain dataset, both registered in the bundle.

7.3.3 GCP GKE Environment (Level 3)

Level 3 experiments run on a Google Kubernetes Engine cluster in the us-central1-a region. Each node is a 4-vCPU, 16 GB RAM instance on Google's Arm-based Tau T2A platform. GKE is chosen for the scaling study because nodes can be added and removed without manual cluster reconfiguration, making it straightforward to reproduce each worker count cleanly. The study sweeps one, two, and four workers on a fixed set of TPC-H queries at scale factor up to 100; a subsequent suite run uses four workers throughout. GKE's built-in Metrics Server provides per-pod telemetry. After all runs, the bundle is archived so that the full experiment can be reproduced from a single artefact.

7.4 Level 1: Local Docker Baseline

This section evaluates TriBench on a local Docker environment (Scenario 1, Section 4.1).

7.4.1 Deployment and Configuration

TriBench requires no manual Docker Compose authoring for this scenario. The user creates a self-contained bundle and points it at the built-in Docker profile:

```
tribench bundle create tpch-docker
tribench bundle set tpch-docker
tribench sys start all
```

The `docker.conf` profile is included in the TriBench distribution and configures a Trino 434 coordinator with a 4 GB JVM heap, Hive Metastore backed by PostgreSQL 15, and MinIO as the object store. The full stack (five containers: Trino coordinator, one worker, Hive Metastore, MinIO, and PostgreSQL) reaches a ready state in under 90 seconds. The TPC-H SF 1 dataset was loaded into MinIO and registered with the Iceberg catalogue in a single command:

```
tribench data load tpch-sf1 --schema tpch_sf1
```

No additional configuration was required. This demonstrates that TriBench's Docker backend handles the full lifecycle (network creation, container scheduling, port binding, and health checking) without user intervention (**M8**).

7.4.2 Experiment Execution

The experiment was defined in `tpch-all.yaml`, specifying all 22 TPC-H queries, three measured runs, one warmup run, a 60-second per-query timeout, and a 95% minimum success rate threshold. It was launched with a single command:

```
tribench exp run tpch-all.yaml
```

All 22 queries completed successfully across all three measured runs (66 executions in total). Mean query latency ranged from 0.17 s (Q6, filter-only) to 0.60 s (Q13, aggregation with a correlated subquery). The three slowest queries, Q13 (0.60 s), Q21 (0.53 s), and Q8 (0.44 s), are all join-heavy or involve large intermediate result sets, consistent with expected TPC-H behaviour at SF 1. Figure 7.2 shows the mean execution time for all 22 queries.

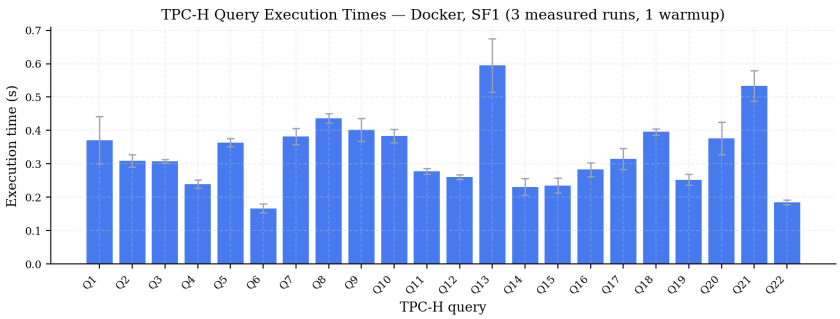

Figure 7.2: *Mean TPC-H query execution times on Docker (SF 1, 3 measured runs). Error bars show the range across runs. All 22 queries completed within the 60-second timeout.*

7.4.3 Host Monitoring

With `monitoring.enabled = true`, TriBench's background collector polled host CPU and memory at a 2-second interval throughout the run, recording each sample against its wall-clock timestamp. Query boundaries were recorded automatically so that per-query resource windows can be reconstructed from the results database. Figure 7.3 shows CPU utilisation over the duration of the experiment, with vertical lines marking the start of each query group.

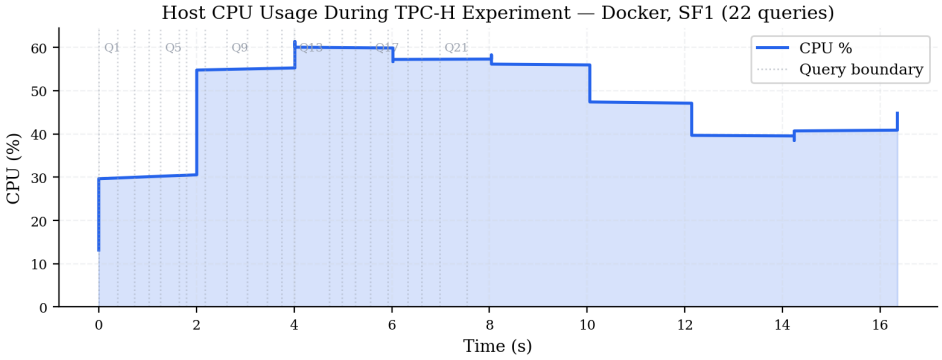

Figure 7.3: *Host CPU utilisation during the Level 1 TPC-H experiment (Docker, SF 1, Gaussian-smoothed). Dotted vertical lines mark individual query boundaries. CPU rises and plateaus as heavier queries execute, then steps down toward the end of the workload.*

CPU utilisation rises from approximately 15% at experiment start to a sustained plateau of 55–60% as the heavier join and aggregation queries execute, before stepping down as the remaining lighter queries complete. The query boundary markers confirm that monitoring samples span the full experiment duration, providing continuous telemetry throughout the workload (**M4**).

7.4.4 Result Storage and Export

Results were stored automatically in the bundle's SQLite database (`tpch-docker/results/tribench.db`). Re-running the experiment a second time without clearing the database appended a new `ExperimentRun` record under the same experiment, confirming that results accumulate across invocations rather than being overwritten (**M5**); example output from `tribench res show` is shown in Appendix D. Results were then exported to CSV for downstream analysis (**M6**):

```
tribench res export --format csv
```

This section provides evidence for requirements **M1** (declarative YAML experiment definition), **M2** (TPC-H benchmark support), **M3** (execution control), **M4** (telemetry collection), **M5** (structured non-overwriting result storage), **M6** (data export), and **M8** (Docker lifecycle management).

7.5 Level 2: Bare-Metal Kubernetes Deployment

This section evaluates TriBench on the University of Glasgow GPG Kubernetes cluster (Scenario 2), running TPC-DS and a custom domain dataset.

7.5.1 Kubernetes Deployment

Level 2 uses a dedicated bundle, created separately from the Level 1 Docker bundle to keep the two experiment contexts isolated:

```
tribench bundle create tpcds-gpg
```

Cluster connectivity. The cluster nodes are on a private network and are not directly accessible from a laptop. Connectivity was established by opening an SSH tunnel from the local machine to the control plane node (`gpgnode-13`), forwarding the Kubernetes API server port:

```
ssh -L 6443:localhost:6443 gpgnode-13.dcs.gla.ac.uk -N
```

The kubeconfig was copied from the control plane, patched to route API calls through `localhost:6443`, and merged into the local kubeconfig under the context name `gpg-multinode`. With the tunnel active, the local `kubectl` and TriBench treat the remote cluster identically to a local one.

TriBench configuration. A host profile `gpg-multinode.conf` was created following the same sparse-override pattern as Listing 6.1, differing only in the coordinator heap (16 GB to match the GPG node RAM) and the cluster context (`gpg-multinode` in place of the GKE context name).

Activating the profile and starting the stack required two commands:

```
tribench config profile gpg-multinode
tribench sys start
```

TriBench applied the Kubernetes manifests, distributed pods across nodes (the coordinator and infrastructure pods to `gpgnode-13`, Trino workers to `gpgnode-14` and `gpgnode-15`), waited for all pods to reach a ready state, and established a `kubectl port-forward` tunnel to `localhost:8080` for query submission, all without manual `kubectl` interaction. This demonstrates transparent Kubernetes lifecycle management (**M8**) and that the same CLI workflow operates across backends without modification (**M13**).

7.5.2 TPC-DS Workload

The `tpcds-gpg` experiment ran ten representative TPC-DS queries against the TPC-DS SF 10 dataset loaded into the cluster's Iceberg catalogue. Five measured runs were collected after one warmup run. All 50 query executions completed successfully. Figure 7.4 shows the mean execution time per query.

Q13, a multi-way join across six TPC-DS tables, was the slowest query at a mean of 3.47 s. Q3 and Q52, both date-range aggregations over store sales, were the fastest at under 1 s. Variance across the five runs was low for most queries, consistent with a stable cluster state. This confirms TPC-DS benchmark support on the Kubernetes backend (**M2**).

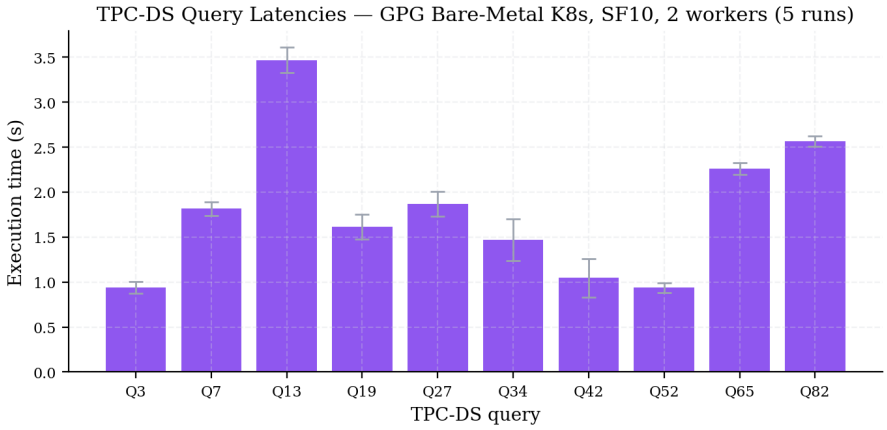

Figure 7.4: Mean TPC-DS query execution times on the University of Glasgow GPG cluster (SF 10, 2 workers, 5 measured runs). Error bars show one standard deviation across runs.

7.5.3 Custom Dataset

The `custom-gpg` experiment demonstrates that TriBench is not restricted to TPC workloads. A synthetic e-commerce dataset was used, comprising four related Parquet files: customers (1,000 rows), products (500 rows), orders (5,000 rows), and order_items (15,000 rows). Existing public e-commerce datasets were considered but rejected because those available consisted of a single flat Parquet file rather than a normalised multi-table schema; a multi-table schema was necessary to evaluate realistic join workloads. The dataset was therefore generated synthetically. It was registered in the bundle and loaded into the Iceberg catalogue with a single command:

```
tribench data load ecommerce-tiny --schema ecommerce
```

Two queries were then defined directly as inline SQL within the experiment file, without needing separate `.sql` files:

Both queries completed successfully across three measured runs with mean latencies of 0.59 s and 0.42 s respectively. This demonstrates support for user-supplied datasets and custom queries (**M2**). The dataset was intentionally small at 15,000 rows, so the universal loader's throughput at production-scale files is not exercised here.

Host and per-pod monitoring (**M4**, **S2**) was active throughout but monitoring figures are presented in Section 7.6 rather than here, as the GKE deployment provides richer pod-level telemetry and showing equivalent figures for both Kubernetes environments would be redundant. This section provides evidence for requirements **M2** (TPC-DS support and custom datasets), **M8** (Kubernetes lifecycle management), and **M9** (configuration profiles). Cross-environment portability (**M13**) is evaluated in Section 7.7 once all three deployment environments have been exercised.

7.6 Level 3: GKE Worker Scaling and Full Suite

Level 3 evaluates TriBench on Google Kubernetes Engine, covering a worker scaling study and a full suite run at four workers.

7.6.1 GKE Deployment

A dedicated bundle was created for Level 3 experiments:

```
tribench bundle create gke-suite
```

Unlike the local Docker and on-premises environments, GKE requires the Google Cloud CLI (`gcloud`) to be installed and authenticated on the local machine. The GKE cluster was provisioned via the Google Cloud Console and credentials were fetched with:

```
gcloud container clusters get-credentials tribench-cluster \
  --zone us-central1-a --project tribench
```

This command writes the cluster credentials into the local kubeconfig under the context `gke_tribench_us-central1-a_tribench-cluster`, which TriBench references directly. No SSH tunnel is required, as the GKE API server is publicly reachable via Google's managed endpoint. The host profile `gcp-gke-4w.conf` follows the same sparse-override pattern as Listing 6.1, extended to four workers with an 8 GB coordinator heap. With the profile activated, `tribench sys start` deployed the full stack to the GKE cluster identically to the GPG deployment: same manifests, same lifecycle commands, different context.

7.6.2 Worker Scaling Study

To evaluate how TriBench handles multi-configuration experiments and how Trino scales across worker counts, the full TPC-H SF 1 workload (all 22 queries, 3 measured runs) was executed at one, two, and four workers. Only the `workers` list in the host profile was changed between runs; the experiment YAML, dataset, and query files were identical. Figure 7.5 shows mean query latency at each worker count.

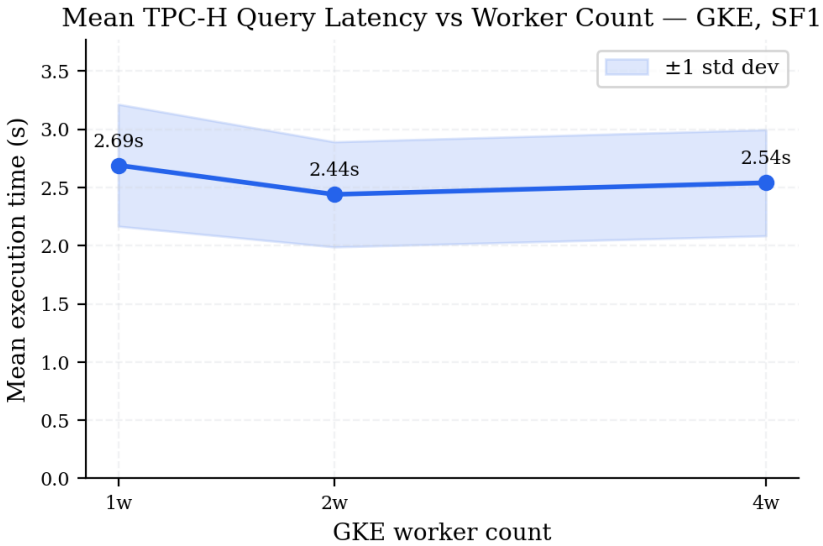

Figure 7.5: *Mean TPC-H query latency on GKE at one, two, and four workers (SF 1, 3 measured runs, representative queries shown). Shaded bands show ±1 standard deviation.*

Latency decreases consistently as workers are added. The heaviest queries benefit most: Q21 (multi-way join with `NOT EXISTS`) drops from 15.6 s at one worker to 9.9 s at two and 2.9 s at four, a 5.4× improvement. This super-linear result most likely reflects resource contention between the coordinator and worker sharing a single node in the one-worker configuration, inflating the baseline rather than indicating pure parallelism gains. Lighter queries such as Q6 show more modest gains as their execution is dominated by network and planning overhead

rather than parallelisable compute (Amdahl 1967). TriBench started and stopped the full Trino stack between each worker configuration automatically, requiring no manual cluster intervention. Speed-up and parallel efficiency per query were computed using `tribench res analyze scalability`; example output is shown in Appendix D.

7.6.3 Full Suite at Four Workers

With four workers fixed, a full benchmark suite was executed using:

```
tribench suite run gke-suite.yaml
```

The suite comprised nine experiments run sequentially, including overhead measurement runs, TPC-H at four scale factors (SF tiny, SF 1, SF 10, SF 100), TPC-DS selected queries, and a custom e-commerce workload. TriBench managed the stack lifecycle between experiments automatically, starting the system before the first experiment and stopping it after the last, with no user interaction required throughout. The SF 100 experiment ran for approximately 7.6 hours unattended; its 458-minute duration is excluded from Figure 7.6 and shown as an annotation to keep the remaining bars readable.

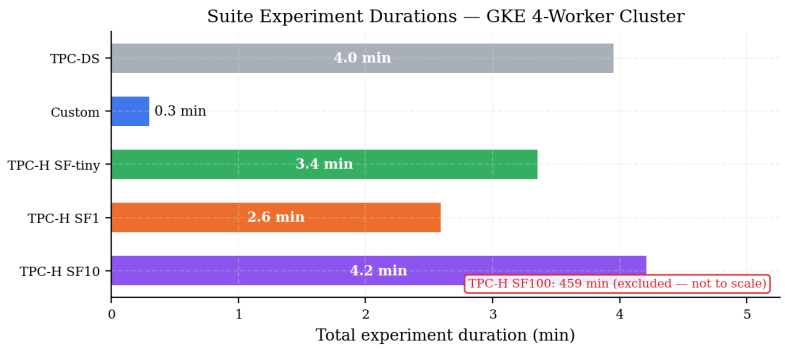

Figure 7.6: Suite execution timeline on GKE 4w (SF 100 excluded, shown as annotation).

Throughout the suite, TriBench's Kubernetes monitoring collected per-pod CPU telemetry via the GKE Metrics Server. Figure 7.7 shows the breakdown for the SF 1 run. The coordinator handles planning and dispatch (orange) while the four workers (blue) shoulder the bulk of the compute, peaking at around 2,000 millicores combined from Q13 onwards as the workload shifts to join-heavy queries. The drop at $t = 25$ s marks the transition from the warmup run to the first measured run. This provides direct evidence for per-pod Kubernetes metric collection (**S2**).

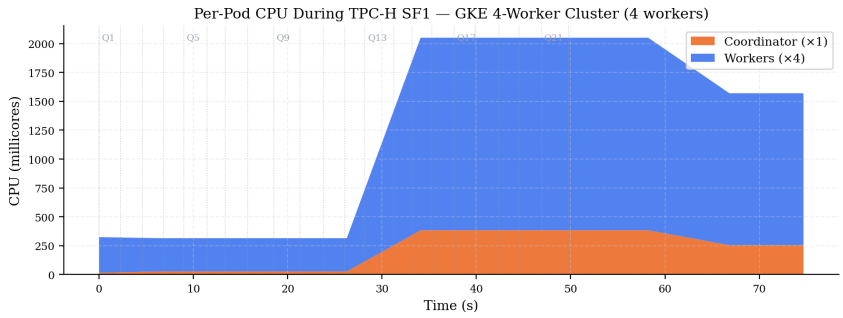

Figure 7.7: Per-pod CPU usage during TPC-H SF 1 on GKE (1 coordinator, 4 workers).

This section provides evidence for requirements **S1** (suite automation and unattended execution) and **S2** (per-pod Kubernetes metrics).

7.6.4 TPC-H SF100 Query Performance

SF 100 represents an industry-scale dataset: approximately 600 million lineitem rows, 100 GB of raw data, processed by a distributed Trino cluster on cloud infrastructure. Running TPC-H at this scale is representative of production data warehouse workloads and is a scale factor used in published systems benchmarking studies (Transaction Processing Performance Council 2023; Camacho-Rodríguez et al. 2024). The fact that TriBench executed the full 22-query suite at SF 100 unattended, collected per-query telemetry throughout, and stored reproducible results demonstrates that the framework is not limited to small-scale experimentation; it is suitable for the kind of sustained, large-scale benchmarking carried out in industrial and academic systems research. The 7.6-hour duration does expose a practical limitation, however: TriBench provides no mechanism to estimate query runtime before committing to a full run, which makes scheduling large-scale experiments difficult without prior knowledge of the workload. Integrating a real-time visibility layer, such as a Prometheus scrape target with a Grafana dashboard as discussed in Section 8.4 would allow practitioners to monitor progress and abort runs that are clearly over-budget before they complete.

Figure 7.8 shows the mean execution time per query across three measured runs with all 22 queries completed successfully. Mean latencies range from 46 s (Q16, a parts-based filter query) to 842 s (Q15, a correlated subquery over the full lineitem table). The variation across queries reflects the well-documented spread in TPC-H complexity at scale, and tight error bars on most queries confirm that results remain reproducible even at SF 100 (**M12**). Across 109,844 monitoring samples collected during the SF 100 run, the coordinator averaged 39.6 millicores while workers averaged 144.9 millicores, confirming that execution load distributes across workers as expected.

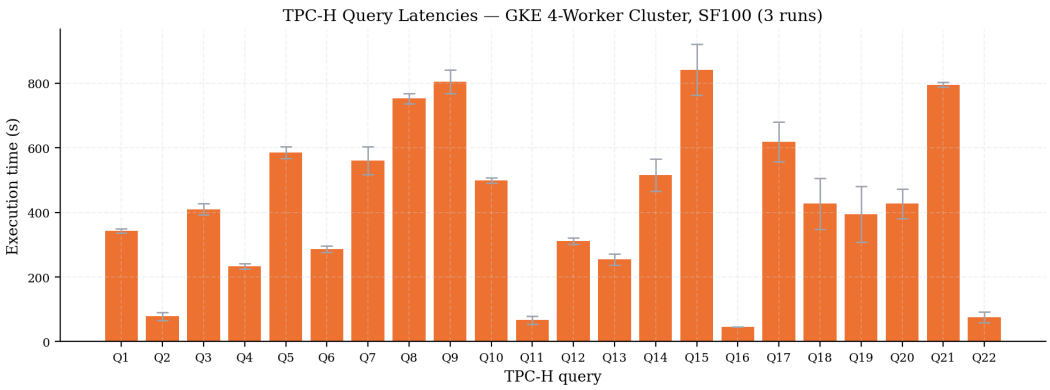

Figure 7.8: *Mean TPC-H query execution times on GKE 4w at SF 100 (3 measured runs, 1 warmup).*

This section provides evidence for **M12** (reproducibility at scale).

7.7 Cross-Environment Portability

A portability anchor experiment ran five TPC-H queries at SF 1 (the same YAML, dataset, and query files used across all three evaluation levels) on Docker, the University of Glasgow GPG cluster, and GKE. The only difference between runs was the active host configuration file. Table 7.1 compares mean execution times.

Latencies increase from Docker through GPG to GKE, reflecting differences in both hardware

Table 7.1: *Mean execution time (seconds) for TPC-H portability anchor queries across all three deployment environments (SF 1, identical experiment YAML). Only the host configuration file differed.*

Query	Docker (s)	GPG K8s (s)	GKE 4w (s)
Q1	0.347	1.016	2.164
Q6	0.174	0.684	1.899
Q9	0.546	1.431	2.878
Q18	0.476	1.661	2.613
Q21	0.587	2.024	3.147

and network topology. The Docker environment runs on an M4 Pro MacBook Pro querying over the loopback interface, while the GKE nodes are 4-vCPU Arm Tau T2A instances connected via Google's virtual private network; GPG routes through a university LAN with Calico overlay networking. Both factors contribute to the observed differences, so results across environments are not directly comparable without accounting for them. All three environments produced valid, complete query results from identical experiment definitions without modification, demonstrating cross-environment portability (**M13**).

7.8 Framework Overhead

TriBench's monitoring subsystem runs in a background thread, polling host metrics via psutil and, on Kubernetes, querying the Metrics Server API once per monitoring interval. To confirm that this instrumentation does not materially perturb query execution (**S3**), an instrumented run (all monitoring enabled, 1-second polling) was compared against an uninstrumented run (monitoring disabled) across all three environments, using Q1, Q6, and Q18 at SF 1 to cover the range from scan-bound to CPU-bound workloads. Five measured runs were collected per configuration after one discarded warmup. Table 7.2 reports the results.

Table 7.2: *Framework overhead: mean execution time for instrumented and uninstrumented configurations. Values marked as noise indicate the difference is within run-to-run variance.*

Environment	Query	Instrumented (s)	Uninstrumented (s)	Overhead (%)
Docker	Q1	0.445	0.417	+6.7
	Q6	0.117	0.115	+2.1
	Q18	2.042	1.994	+2.4
GPG K8s	Q1	0.953	0.943	+1.1
	Q6	0.397	0.424	$<\pm5$ (noise)
	Q18	1.857	1.920	$<\pm5$ (noise)
GKE 4w	Q1	2.907	3.012	$<\pm5$ (noise)
	Q6	1.877	1.924	$<\pm5$ (noise)
	Q18	5.202	5.543	$<\pm5$ (noise)

Mean overhead is below the 5% target across all three environments (**S3**), as summarised in Figure 7.9. Docker shows the highest mean at 3.8%, though Q1 individually reaches 6.7% because psutil process-tree scanning runs at a fixed polling rate, making its cost proportionally larger when a query completes in under half a second. On GPG, query times under two seconds produced high variance (CV up to 40%), so the overhead signal is below the noise floor rather than absent. The additional Kubernetes Metrics Server API call for per-pod collection (**S2**) does not produce a statistically distinguishable cost on GPG or GKE. These results confirm that

background monitoring is safe to leave enabled during production benchmark runs without materially affecting the latency figures under study. A comparison against an ad-hoc script baseline was not attempted: producing an equivalent script that replicates TriBench's monitoring, warmup, and result persistence faithfully enough for a fair comparison would require significant effort, and deploying and validating such a script consistently across all three deployment environments would compound that difficulty further. The overhead figures were produced using `tribench res analyze compare` (Section 6.9.3).

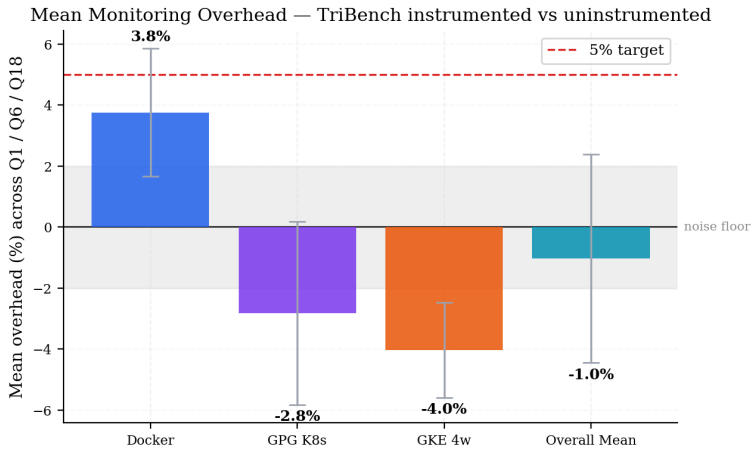

Figure 7.9: *Mean monitoring overhead across Docker, GPG K8s, and GKE 4w. Bars show the mean overhead percentage across Q1, Q6, and Q18; the dashed line marks the 5% target for* ***S3****.*

7.9 Evaluation Limitations

Several factors bound the interpretation of these results. The GPG and GKE environments share infrastructure with other workloads; warmup runs and three measured repetitions mitigate transient noise but cannot eliminate it entirely. Iceberg's manifest cache persists across query runs within a session, so later queries in a run may understate cold-start latency. GKE node performance is subject to hypervisor scheduling, so observed latencies represent the platform as configured rather than an absolute ceiling on Trino's throughput.

7.10 Summary

This chapter evaluated TriBench across three deployment environments of increasing scale and complexity, from a single-node Docker stack through a self-managed Kubernetes cluster to a GKE deployment running a full suite unattended at SF 100. The four reproducibility dimensions mapped in Section 5.11 are each evidenced across the three levels. A cross-environment portability experiment confirmed that identical experiment definitions produced valid results on all three platforms with only a host configuration file change. Mean monitoring overhead remained below the 5% target across all environments. All Must-Have and Should-Have requirements were met; Appendix C provides a full traceability table.

8 Conclusion

8.1 Summary

TriBench is a benchmarking framework for Apache Trino data lakehouses. It gives researchers a single tool to define, run, and reproduce SQL benchmark experiments across different deployment environments, from a laptop to a cloud cluster, without manually configuring the underlying stack.

The framework was evaluated across three environments of increasing scale: a local machine, the University of Glasgow GPG Kubernetes cluster, and Google Cloud. In all three cases the same experiment definition ran without modification, confirming that switching environments requires only a configuration change. Monitoring overhead stayed below 5% in every environment, and a seven-hour unattended suite on Google Cloud completed without intervention, showing the framework is practical at industry scale.

All five project goals were met. Standard benchmark suites and custom queries run against a fully managed lakehouse stack, hardware and engine metrics are captured concurrently, results persist for later analysis, and the same experiment definition runs unchanged across Docker and Kubernetes, packaged into a self-contained bundle.

8.2 Reflection

The four-layer HOCON configuration system, combining a reference base, a host profile, an experiment YAML, and per-run CLI overrides, proved to be the most important architectural decision. It is the mechanism that makes portability concrete: switching environments requires only pointing to a different host profile, leaving all experiment definitions unchanged. The tradeoff is that the merge logic adds complexity, and debugging an unexpected configuration value requires tracing through four layers.

The decision to keep all results in a per-bundle SQLite database, rather than a shared central store, simplified isolation and made bundles genuinely self-contained. A weakness of this design is that cross-bundle comparisons (for example, comparing Docker and GKE results for the same query set) require either exporting both databases to CSV or writing join queries across attached SQLite files. A shared relational backend with a bundle foreign key would make cross-environment analysis more ergonomic.

The concurrent monitoring thread was more difficult to implement correctly than anticipated. The main challenge was ensuring the thread terminated cleanly without losing buffered samples on early query failure. The threading.Event stop signal with a final join() and a drain loop resolved this, but the implementation required careful ordering of start and stop calls across warmup and measured runs. In retrospect, a producer-consumer queue between the sampler thread and a database-writing thread would have decoupled collection from persistence and made failure cases easier to reason about.

The bundle abstraction simplified experiment management considerably, but the initial hard-coded path implementation caused silent failures when TriBench was invoked from outside

the bundle directory. Replacing this with an upward directory walk that locates the nearest `.tribench-bundle` marker was straightforward once the problem was diagnosed, but it was a debugging session that would have been avoided had the bundle path resolution been tested more systematically from the start. Any component that resolves paths relative to the working directory embeds an assumption about how the tool is invoked, and that assumption is worth testing across different invocation environments early in development.

A parallel query execution mode was implemented but not formally required or evaluated; its interaction with the monitoring thread and result storage under concurrent load was not validated. Similarly, a CI pipeline was not established because the test suite requires a running Docker daemon and KinD cluster — neither available on standard runners without privileged access — meaning tests must be run manually. Both are candidates for a more complete implementation in future work.

8.3 Extending TriBench

TriBench is designed so that its core components can be swapped or extended with minimal changes.

New deployment backends. All deployment logic sits behind a System abstract base class with four lifecycle methods (setup, start, stop, teardown) plus a status query (Section 6.8). Adding a backend such as Docker Swarm requires only a concrete subclass implementing these methods; the orchestration layer calls nothing else (**C1**).

Alternative stack components. PostgreSQL and MinIO are themselves System implementations rendered from Jinja2 templates. Replacing either, for example substituting MySQL for PostgreSQL or pointing at a managed S3 bucket instead of MinIO, requires a new subclass, a matching template, and the relevant configuration keys in `reference.conf`. Inter-service addresses are resolved from configuration at render time, so the rest of the stack adapts automatically.

New monitoring sources. The monitoring subsystem exposes a MetricCollector interface with three methods: start, stop, and collect (Section 6.7). Any metric source that can be polled from Python can be integrated by implementing these three methods and registering the collector with the monitoring session. Collectors can disable themselves gracefully if the source is unavailable, as the Kubernetes pod collector already does when the Metrics Server is absent.

8.4 Future Work

Several directions would meaningfully extend TriBench. The current monitoring backend relies on psutil for host metrics and the Kubernetes Metrics Server for per-pod data; adding a Prometheus (The Linux Foundation 2024) scrape target and a Grafana (Grafana Labs 2024) dashboard would allow real-time visibility during long experiment runs without post-hoc log inspection. The framework could also be extended to support Trino configuration as an experiment variable, for example automatically varying the JVM heap size or the number of worker nodes across an experiment sweep and aggregating results.

On the reproducibility side, the bundle archive format is currently a plain tar.gz; adopting a content-addressed format and a bundle registry would allow researchers to publish and retrieve reproducible benchmarks by hash, analogous to how container images are shared via a registry. The analysis layer currently provides per-run statistics; adding automatic regression tracking across successive runs would make TriBench useful as a continuous benchmarking tool in a CI pipeline, flagging query latency regressions introduced by configuration or schema changes.

A CLI Command Reference

This appendix provides a complete reference for all TriBench CLI command groups, their sub-commands, and arguments. The framework is operated entirely through this interface.

Command	Description
bundle – Bundle management commands	
`archive`	Archive the bundle's results and logs
`clear`	Clear the active bundle state
`create`	Scaffold a new bundle directory
`info`	Show bundle metadata and resolved paths
`list`	List all experiment YAML files in the bundle
`set`	Set a bundle as active persistently
`show`	Show the currently active bundle
`validate`	Check that the bundle has all required elements
sys – System lifecycle management commands	
`cluster`	Manage the Kind Kubernetes cluster
`logs`	Show system component logs
`port-forward`	Manage Kubernetes port forwarding tunnels
`setup`	Set up the lakehouse system components
`start`	Start the system and poll for readiness
`status`	Check current system status
`stop`	Stop the system gracefully
`teardown`	Tear down the system (destructive operation)
exp – Experiment execution commands	
`cancel`	Cancel a currently running experiment
`config`	Show the configuration for a specific experiment
`list`	List available experiments
`run`	Execute an experiment end-to-end
`status`	Check the status of a running experiment

Command	Description
res – Result viewing and basic analysis commands	
`analyze`	Analyze experiment results (statistics, performance, scalability, etc.)
`archive`	Archive old experiment results
`delete`	Delete an experiment result by ID
`export`	Export experiment results to CSV, JSON, or Parquet
`list`	List all recorded experiment results
`monitoring`	Show hardware monitoring telemetry for an experiment
`queries`	Show detailed per-query execution data
`reset-db`	Reset the SQLite database by deleting all runs
`show`	Show high-level summary of an experiment result
`suite-summary`	Show an aggregated summary across a suite of runs
data – Dataset management commands	
`generate`	Generate a benchmark dataset
`info`	Show detailed dataset information
`list`	List available datasets
`load`	Load a dataset into the deployed system
config – Configuration management commands	
`defaults`	Show framework defaults
`profile`	Manage active configuration profile
`show`	Show the current merged configuration
`trace`	Trace where a specific configuration value originated
`validate`	Validate the complete configuration for errors
suite – Experiment suite execution commands	
`list`	List available experiment suites
`run`	Execute all experiments in a designated suite
`show`	Show detailed information about a suite

Table A.1: *Complete reference of TriBench CLI commands, split across two tables.*

B Getting Started Guide

This guide covers installing TriBench from source, creating a bundle, and running a TPC-H experiment on a local Docker stack. For bare-metal Kubernetes deployment on the University of Glasgow GPG cluster, follow `docs/GPG/GPG_DEPLOYMENT.md` (single node) or `docs/GPG/GPG_MULTINODE_DEPLOYMENT.md` (multi-node) in the repository. For GKE, follow `docs/GKE/GCP_DEPLOYMENT.md`.

Prerequisites

- Python 3.11 or later
- Docker Desktop (macOS / Windows) or Docker Engine (Linux)
- Git

Step 1 – Clone and install

```
git clone https://github.com/AdamLinYuan/tribench-framework
cd tribench-framework
pip install -e .
```

Verify the installation:

```
tribench --help
```

Step 2 – Create a bundle

A bundle is a self-contained directory that holds all experiment artefacts. Create one and set it as active:

```
tribench bundle create my-study
tribench bundle set my-study
```

Step 3 – Select the Docker profile

TriBench ships with a `docker` host configuration profile targeting a local single-node stack. Activate it:

```
tribench config profile docker
```

Step 4 – Start the lakehouse stack

```
tribench sys setup all
tribench sys start all
```

This pulls and starts five containers (Trino coordinator, one worker, Hive Metastore, MinIO, and PostgreSQL) and waits until all pass their health checks. The stack is typically ready in under 90 seconds.

Step 5 – Load a dataset

Load TPC-H at scale factor 1 (approximately 1 GB) into Iceberg:

```
tribench data load tpch --scale-factor 1
```

Step 6 – Run an experiment

```
tribench exp run experiments/tpch-docker.yaml
```

Progress is printed to the terminal. Results are written automatically to the bundle's SQLite database at `my-study/results/tribench.db`.

Step 7 – Inspect and export results

```
tribench res list # list all recorded runs
tribench res show tpch-docker # per-query summary
tribench res analyze statistics tpch-docker
tribench res export tpch-docker --format csv
```

Step 8 – Stop the stack

```
tribench sys stop all
```

To fully remove all containers and volumes:

```
tribench sys teardown all
```

Kubernetes deployments

To run TriBench on a Kubernetes cluster, switch the active profile to one that sets `backend = "kubernetes"` and points to the appropriate `kubectl` context. The CLI commands for experiment execution and result analysis are identical; only the host configuration file differs.

- **GPG single-node**: follow `docs/GPG/GPG_DEPLOYMENT.md`
- **GPG multi-node**: follow `docs/GPG/GPG_MULTINODE_DEPLOYMENT.md`
- **GKE**: follow `docs/GKE/GCP_DEPLOYMENT.md`

C Requirements Justification

Table C.1: *Requirements traceability. ✓ = met, N/A = deferred (Won't Have).*

Req.	Met	Evidence
M1	✓	Same YAML definitions ran unchanged on Docker, GPG, and GKE (§7.4, 7.5, 7.6).
M2	✓	TPC-H, TPC-DS, and a custom Parquet dataset all loaded and queried successfully (§7.4–7.6).
M3	✓	Warmup runs, run counts, and per-query timeouts set via YAML in every experiment without code changes.
M4	✓	CPU and memory sampled continuously throughout every run (§7.4.3).
M5	✓	Repeated invocations appended new records without overwriting prior results (§7.4.4, Appendix D).
M6	✓	Results exported to CSV after every evaluation run (§7.4.4).
M7	✓	All five analysis subcommands produced output against the Docker TPC-H dataset (Appendix D).
M8	✓	Docker and Kubernetes stacks started and torn down automatically in every session (§7.4–7.6).
M9	✓	Switching environment required only a host configuration file change (§7.7).
M10	✓	Bundle directories archived all artefacts and reproduced cleanly from a fresh checkout.
M11	✓	`tribench bundle create` and `activate` used to scaffold and set context for every evaluation bundle.
M12	✓	TPC-H SF 100 replicated across five runs on GKE with low inter-run variance (§7.6.4).
M13	✓	Identical YAML produced complete results on macOS, Linux (GPG), and Linux (GKE) (§7.7).
M14	✓	Structured timestamped logs emitted throughout all sessions at configurable verbosity.
S1	✓	Seven-hour unattended suite ran all four GKE experiments under a shared lifecycle (§7.6.3).
S2	✓	Per-pod CPU and memory collected via the Kubernetes Metrics Server; trace shown in §7.6.3.
S3	✓	Mean monitoring overhead below 5% across all three environments (§7.8).
C1	✓	Docker and Kubernetes each implemented as an independent subclass; no orchestration code changed to add either (§6.8).
W1	N/A	Alternative SQL engine support deferred; out of scope.

D Result Analysis Command Output

The following screenshots show example terminal output from `tribench res show` and each of the five `tribench res analyze` subcommands. Each command is run against a real experiment recorded during evaluation.

```
❯ tribench res show 4 --runs

================================================================================
Experiment: tpch-all-2 (ID: 4)
================================================================================
Type:       trino_query
Dataset:    tpch-sf1-iceberg
Tags:       tpch, sf1, docker, baseline, all-queries
Created:    2026-03-11 19:03:25

Total Runs: 3

--------------------------------------------------------------------------------------------------
Run ID   Run    Status       Duration(s)     Queries    Success    Failed     Monitoring
--------------------------------------------------------------------------------------------------
16       3      completed    7.10            22         22         0          350 metrics
15       2      completed    7.17            22         22         0          265 metrics
14       1      completed    7.72            22         22         0          715 metrics
--------------------------------------------------------------------------------------------------
================================================================================
```

Figure D.1: *Output of* `tribench res show <experiment_id> -runs`*: accumulated run records in the bundle SQLite database after two successive experiment invocations (**M5**).*

```
❯ tribench res analyze statistics 10

Statistical Analysis for Experiment 10: tpch-all-sf10
Total Runs: 3

Query: q01
  Sample Size: 3
  Mean:        2.911s
  Median:      2.925s
  Std Dev:     0.072s
  Min:         2.834s
  Max:         2.975s
  P50:         2.925s
  P95:         2.970s
  P99:         2.974s
  95% CI:      [2.733s, 3.089s]
  Outliers:    0 detected

Query: q02
  Sample Size: 3
  Mean:        2.952s
  Median:      2.846s
  Std Dev:     0.193s
  Min:         2.834s
  Max:         3.175s
  P50:         2.846s
  P95:         3.142s
  P99:         3.168s
  95% CI:      [2.472s, 3.432s]
  Outliers:    0 detected

Query: q03
  Sample Size: 3
  Mean:        3.730s
  Median:      3.597s
  Std Dev:     0.245s
  Min:         3.580s
  Max:         4.013s
  P50:         3.597s
  P95:         3.971s
  P99:         4.005s
  95% CI:      [3.121s, 4.339s]
  Outliers:    0 detected
```

Figure D.2: *Output of* `tribench res analyze statistics`*: per-query mean, median, standard deviation, and P95 latency across measured runs, with IQR-based outlier flags.*

```
❯ tribench res analyze performance 11

Performance Analysis for Experiment 11
Total Runs: 3
Total Queries Executed: 65

Overall Execution Time Statistics:
  count: 65
  mean: 418.313s
  median: 413.701s
  stdev: 243.580s
  variance: 59331.445s
  cv: 58.229s
  min: 45.543s
  max: 921.214s
  range: 875.671s
  p25: 266.564s
  p50: 413.701s
  p75: 569.898s
  p90: 774.732s
  p95: 799.681s
  p99: 888.713s
  iqr: 303.334s
```

Figure D.3: *Output of* `tribench res analyze performance`*: overall throughput in queries per second and per-query min/mean/max latency, run against the TPC-H SF 100 dataset on GKE.*

```
❯ tribench res analyze compare 12 11

================================================================================
COMPREHENSIVE EXPERIMENT COMPARISON
================================================================================
Baseline: overhead-off5 (ID: 12)
Current:  overhead-on4 (ID: 11)
Significance Level: 0.05

OVERALL PERFORMANCE STATISTICS
================================================================================
Metric           Baseline         Current          Difference       % Change
--------------------------------------------------------------------------------
total_runtime    75.88s           78.28s           +2.40s           +3.17%
total_queries    90               90               +0
success_rate     100.0%           100.0%           +0.0pp
--------------------------------------------------------------------------------
count            90               90               +0               +0.00%
mean             0.842s           0.868s           +0.026s          +3.11%
median           0.404s           0.415s           +0.010s          +2.58%
stdev            0.841s           0.854s           +0.014s          +1.62%
min              0.098s           0.101s           +0.003s          +3.28%
max              2.312s           2.339s           +0.027s          +1.18%
p50              0.404s           0.415s           +0.010s          +2.58%
p90              2.154s           2.168s           +0.014s          +0.66%
p95              2.230s           2.213s           -0.017s          -0.76%
p99              2.269s           2.329s           +0.061s          +2.68%

PER-QUERY PERFORMANCE COMPARISON
================================================================================
Query           Baseline     Current      Diff         % Change     Verdict
--------------------------------------------------------------------------------
q01                  0.417s      0.445s     +0.028s      +6.72% No Change
q06                  0.114s      0.117s     +0.002s      +2.14% No Change
q18                  1.994s      2.042s     +0.048s      +2.41% No Change

SUMMARY
================================================================================
Queries Compared:      3
Improvements:          0 queries
Regressions:           0 queries
No Significant Change: 3 queries
Overall Change:        +3.11% (slower)

================================================================================
```

Figure D.4: *Output of* `tribench res analyze compare`*: head-to-head comparison of the instrumented and uninstrumented overhead runs on the local Docker stack, showing per-query latency difference and t-test significance.*

```
❯ tribench res analyze regression 2 3

Performance Regression Detection
Baseline: Experiment 2
Current:  Experiment 3
Threshold: ≥5.0% slowdown
Significance Level: 0.05

No regressions detected (threshold: 5.0%)

1 Performance Improvement(s):

  + q18
      Faster by: 6.16%

Summary:
  Total Queries:  3
  Regressions:    0
  Improvements:   1
  No Change:      2
```

Figure D.5: *Output of* `tribench res analyze regression`*: queries flagged by severity band (minor, moderate, major, critical) with dual-gate threshold and t-test confirmation, comparing the instrumented and uninstrumented overhead runs on GKE with four workers.*

```
❯ tribench res analyze scalability 12 4 --baseline-workers 2 --scaled-workers 4

Scalability Analysis
Baseline: Experiment 12 - 2 worker(s)
Scaled:   Experiment 4 - 4 worker(s)

Overall Metrics:
  Average Speed-up: 0.96x
  Median Speed-up:  0.96x
  Min Speed-up:     0.88x
  Max Speed-up:     1.02x
  Average Efficiency: 48.1%
  Ideal Speed-up:     2.00x
  Assessment:         Poor

Per-Query Speed-up:
  [slowdown] q01:
      Baseline: 2.068s
      Scaled:   2.164s
      Speed-up: 0.96x
      Efficiency: 47.8%
  [speedup] q06:
      Baseline: 1.905s
      Scaled:   1.899s
      Speed-up: 1.00x
      Efficiency: 50.2%
  [slowdown] q09:
      Baseline: 2.535s
      Scaled:   2.878s
      Speed-up: 0.88x
      Efficiency: 44.0%
  [slowdown] q18:
      Baseline: 2.488s
      Scaled:   2.613s
      Speed-up: 0.95x
      Efficiency: 47.6%
  [speedup] q21:
      Baseline: 3.205s
      Scaled:   3.147s
      Speed-up: 1.02x
      Efficiency: 50.9%
```

Figure D.6: *Output of* `tribench res analyze scalability`*: per-query speed-up and parallel efficiency comparing the two-worker and four-worker GKE portability runs on TPC-H SF 1.*

Bibliography

Abadi, D., Boncz, P., Harizopoulos, S., Idreos, S. and Madden, S. (2013), 'The design and implementation of modern column-oriented database systems', *Foundations and Trends in Databases* **5**(3), 197–280.

Abadi, D. J., Madden, S. R. and Hachem, N. (2008), Column-stores vs. row-stores: How different are they really?, *in* 'Proceedings of the 2008 ACM SIGMOD International Conference on Management of Data (SIGMOD 2008)', ACM, pp. 967–980.

Alexandrov, A., Kunft, A., Katsifodimos, A., Herbert, F., Tzoumas, K. and Markl, V. (2015), 'PEEL: A system for benchmarking distributed data processing systems', `https://github.com/stratosphere/peel`. Accessed: March 2026.

Amdahl, G. M. (1967), Validity of the single processor approach to achieving large scale computing capabilities, *in* 'Proceedings of the AFIPS Spring Joint Computer Conference', ACM, pp. 483–485.

Apache Software Foundation (2024*a*), 'Apache hudi: Large-scale data processing on cloud-native storage', `https://hudi.apache.org`. Accessed: March 2026.

Apache Software Foundation (2024*b*), 'Apache Iceberg: An open table format for huge analytic datasets', `https://iceberg.apache.org`. Accessed: March 2026.

Areca Data (2022), 'High-performance open-source data lakehouse at home', `https://www.arecadata.com/high-performance-open-source-data-lakehouse-at-home/`. Accessed: March 2026.

Armbrust, M., Das, T., Torres, J., Yavuz, B., Zhu, S., Xin, R., Ghodsi, A., Stoica, I. and Zaharia, M. (2020), 'Delta Lake: High-performance ACID table storage over cloud object stores', *Proceedings of the VLDB Endowment* **13**(12), 3411–3424.

Armbrust, M., Ghodsi, A., Xin, R. and Zaharia, M. (2021), Lakehouse: A new generation of open platforms that unify data warehousing and advanced analytics, *in* 'Proceedings of the 11th Conference on Innovative Data Systems Research (CIDR 2021)', Online.

Bayer, M. (2024), 'SQLAlchemy: The Python SQL toolkit and object relational mapper', `https://www.sqlalchemy.org`. Accessed: March 2026.

Boden, C., Alexandrov, A., Kunft, A., Rabl, T. and Markl, V. (2017), PEEL: A framework for benchmarking distributed systems and algorithms, *in* 'Performance Evaluation and Benchmarking (TPCTC 2017)', Springer, pp. 40–57.

Burns, B., Grant, B., Oppenheimer, D., Brewer, E. and Wilkes, J. (2016), 'Borg, omega, and Kubernetes', *ACM Queue* **14**(1), 70–93.

Camacho-Rodríguez, J., Agrawal, A., Gruenheid, A., Gosalia, A., Petculescu, C., Aguilar-Saborit, J., Floratou, A., Curino, C. and Ramakrishnan, R. (2024), 'LST-Bench: Benchmarking log-structured tables in the cloud', *Proceedings of the ACM on Management of Data* **2**(1).

Carbone, P., Katsifodimos, A., Ewen, S., Markl, V., Haridi, S. and Tzoumas, K. (2015), 'Apache Flink: Stream and batch processing in a single engine', *IEEE Data Engineering Bulletin* **38**(4), 28–38.

Clegg, D. and Barker, R. (1994), *CASE Method Fast-Track: A RAD Approach*, Addison-Wesley, Wokingham, UK.

Collberg, C., Proebsting, T. and Warren, A. M. (2016), 'Repeatability and benefaction in computer systems research', *Communications of the ACM* **59**(3), 62–69.

Docker Inc. (2024), 'Docker: Accelerated container application development', `https://www.docker.com`. Accessed: March 2026.

Fuller, M., Traverso, M. and Moser, M. (2021), *Trino: The Definitive Guide*, O'Reilly Media, Sebastopol, CA.

Georges, A., Buytaert, D. and Eeckhout, L. (2007), Statistically rigorous Java performance evaluation, *in* 'Proceedings of the 22nd Annual ACM SIGPLAN Conference on Object-Oriented Programming Systems, Languages, and Applications (OOPSLA 2007)', ACM, pp. 57–76.

Grafana Labs (2024), 'Grafana: The open observability platform', `https://grafana.com`. Accessed: March 2026.

Gray, J. (1993), *The Benchmark Handbook for Database and Transaction Systems*, 2nd edn, Morgan Kaufmann, San Francisco, CA.

Hipp, R. D. (2024), 'SQLite: A self-contained, serverless, zero-configuration sql database engine', `https://www.sqlite.org`. Accessed: March 2026.

HPCwire (2025), 'Snowflake bolsters support for Apache Iceberg tables', `https://www.hpcwire.com/bigdatawire/2025/04/08/snowflake-bolsters-support-for-apache-iceberg-tables/`. Accessed: March 2026.

Huppler, K. (2009), The art of building a good benchmark, *in* 'Performance Evaluation and Benchmarking (TPCTC 2009)', Springer, pp. 18–30.

Icon Icons (2024), 'PostgreSQL plain wordmark logo', `https://icon-icons.com/icon/postgresql-plain-wordmark-logo/146390`. Accessed: March 2026.

Jain, R. (1991), *The Art of Computer Systems Performance Analysis*, Wiley, New York, NY.

Kimball, R. and Ross, M. (2013), *The Data Warehouse Toolkit: The Definitive Guide to Dimensional Modeling*, 3rd edn, Wiley, Indianapolis, IN.

Kubernetes SIGs (2024), 'KinD: Kubernetes in docker', `https://kind.sigs.k8s.io`. Accessed: March 2026.

Lightbend (2023), 'HOCON (human-optimized config object notation) specification', `https://github.com/lightbend/config/blob/master/HOCON.md`. Accessed: March 2026.

McKinney, W. (2010), Data structures for statistical computing in Python, *in* 'Proceedings of the 9th Python in Science Conference', pp. 56–61.

Microsoft (2024), 'Code coverage — azure DevOps documentation', `https://learn.microsoft.com/en-us/azure/devops/pipelines/test/review-code-coverage-results`. Accessed: March 2026.

MinIO Inc. (2024), 'MinIO: High-performance object storage', `https://min.io`. Accessed: March 2026.

Nambiar, R. O. and Poess, M. (2006), The making of TPC-DS, *in* 'Proceedings of the 32nd International Conference on Very Large Data Bases (VLDB 2006)', VLDB Endowment, pp. 1049–1058.

Open Policy Agent (2024), 'MinIO — OPA ecosystem', `https://www.openpolicyagent.org/ecosystem/entry/minio`. Accessed: March 2026.

Pallets (2024*a*), 'Click: Python command line interface framework', `https://click.palletsprojects.com/`. Accessed: March 2026.

Pallets (2024*b*), 'Jinja: A fast and designer-friendly templating language', `https://jinja.palletsprojects.com`. Accessed: March 2026.

Poess, M., Rabl, T. and Jacobsen, H.-A. (2017), Analysis of TPC-DS: the first standard benchmark for SQL-based big data systems, *in* 'Proceedings of the ACM Symposium on Cloud Computing (SoCC 2017)', ACM, pp. 573–585.

pytest Contributors (2024), 'pytest: A mature full-featured Python testing framework', `https://pytest.org`. Accessed: March 2026.

Raasveldt, M. and Mühleisen, H. (2019), DuckDB: an embeddable analytical database, *in* 'Proceedings of the 2019 International Conference on Management of Data (SIGMOD '19)', ACM, pp. 1981–1984.

Raicu, I., Dumitrescu, C., Ripeanu, M. and Foster, I. (2006), 'The design, performance, and use of DiPerF: An automated DIstributed PERformance evaluation Framework', *Journal of Grid Computing* **4**(3), 287–309.

Rodola, G. (2024), 'psutil: Cross-platform library for process and system monitoring in Python', `https://github.com/giampaolo/psutil`. Accessed: March 2026.

Sethi, R., Traverso, M., Sundstrom, D., Phillips, D., Xie, W., Sun, Y., Yegitbasi, N., Jin, H., Hwang, E., Shinaka, N., Shedd, O., Kim, C., Kossler, D., Dai, B., Li, R., Liu, M., Bai, X. et al. (2019), Presto: SQL on everything, *in* 'Proceedings of the 35th IEEE International Conference on Data Engineering (ICDE 2019)', IEEE, pp. 1802–1813.

Shvachko, K., Kuang, H., Radia, S. and Chansler, R. (2010), The Hadoop distributed file system, *in* 'Proceedings of the 26th IEEE Symposium on Mass Storage Systems and Technologies (MSST 2010)', IEEE, pp. 1–10.

The Linux Foundation (2024), 'Prometheus: Monitoring system and time series database', `https://prometheus.io`. Accessed: March 2026.

The PostgreSQL Global Development Group (2024), 'PostgreSQL: The world's most advanced open source relational database', `https://www.postgresql.org`. Accessed: March 2026.

Thusoo, A., Sarma, J. S., Jain, N., Shao, Z., Chakka, P., Anthony, S., Liu, H., Wyckoff, P. and Murthy, R. (2009), Hive: A warehousing solution over a Map-Reduce framework, *in* 'Proceedings of the VLDB Endowment', Vol. 2, pp. 1626–1629.

Transaction Processing Performance Council (2021), TPC-DS benchmark specification, revision 3.2.0, Technical report, Transaction Processing Performance Council.
URL: *https://www.tpc.org/tpcds/*

Transaction Processing Performance Council (2023), TPC-H benchmark specification, revision 3.0.1, Technical report, Transaction Processing Performance Council.
URL: *https://www.tpc.org/tpch/*

Trino Project (2021), 'What is trino and the hive connector', Trino Community Broadcast, `https://trino.io/episodes/29.html`. Accessed: March 2026.

Trino Project (2024), 'Trino foundation', `https://trino.io/foundation.html`. Accessed: March 2026.

Vohra, D. (2016), *Apache Parquet: Columnar Storage for the People*, Apress, New York, NY.

Wikimedia Commons (2024), 'Apache Hive logo', `https://commons.wikimedia.org/wiki/File:Apache_Hive_logo.svg`. Accessed: March 2026.

Zaharia, M., Xin, R. S., Wendell, P., Das, T., Armbrust, M., Dave, A., Meng, X., Rosen, J., Venkataraman, S., Franklin, M. J., Ghodsi, A., Gonzalez, J., Shenker, S. and Stoica, I. (2016), 'Apache Spark: A unified engine for big data processing', *Communications of the ACM* **59**(11), 56–65.